

https://svs.gsfc.nasa.gov/vis/a010000/a011000/a011003/DynamicEarth-Still4_03561.jpg

Robotics for Future Industrial Applications

Learning Global and Local Dynamic Models


Philipp Ennen, M.Sc.



$$p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t) = \mathcal{N}(f(\mathbf{x}_t, \mathbf{u}_t), \Sigma)$$


$$f(\mathbf{x}_t, \mathbf{u}_t) \approx \mathbf{A}_t \mathbf{x}_t + \mathbf{B}_t \mathbf{u}_t$$

$$\mathbf{A}_t = \frac{\partial f}{\partial \mathbf{x}_t} \quad \mathbf{B}_t = \frac{\partial f}{\partial \mathbf{u}_t}$$



What kind of models can we use?

Gaussian process



GP with input (\mathbf{x}, \mathbf{u}) and output \mathbf{x}'

Pro: very data-efficient

Con: not great with non-smooth dynamics

Con: very slow when dataset is big

Neural Network




image: Punjani & Abbeel '14


Input is (\mathbf{x}, \mathbf{u}) , output is \mathbf{x}'

Pro: very expressive, can use lots of data

Con: not so great in low data regimes

This week's focus!

Gaussian Mixture Model




GMM over $(\mathbf{x}, \mathbf{u}, \mathbf{x}')$ tuples

Train on $(\mathbf{x}, \mathbf{u}, \mathbf{x}')$, condition to get $p(\mathbf{x}'|\mathbf{x}, \mathbf{u})$

For i 'th mixture element, $p_i(\mathbf{x}, \mathbf{u})$ gives region where the mode $p_i(\mathbf{x}'|\mathbf{x}, \mathbf{u})$ holds

Pro: very expressive, if the dynamics can be assumed as piecewise linear



1. Run time-varying policy $q(\mathbf{u}_t|\mathbf{x}_t)$ on robot N times
2. Collect dataset $\mathcal{D} = \{\tau_i\}$ where $\tau_i = \{\mathbf{x}_{1i}, \mathbf{u}_{1i}, \dots, \mathbf{x}_{Ti}, \mathbf{u}_{Ti}\}$
3. For each $t \in \{0, \dots, T-1\}$, fit linear Gaussian $p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t)$
4. Solve control problem to get new $q(\mathbf{u}_t|\mathbf{x}_t)$

I. Gaussians

- I. Univariate Gaussian
- II. Multivariate Gaussian
- III. Conditioning (Bayes' rule)

II. Gaussian Mixture Model

III. Learning Local Dynamic Models


Disclaimer: lots of linear algebra now. In fact, pretty much all computations with Gaussians will be reduced to linear algebra!

Univariate Gaussian

- Gaussian distribution with mean μ , and standard deviation σ :

$$X \sim \mathcal{N}(\mu, \sigma^2)$$

$$p(x; \mu, \sigma^2) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$



Properties of Gaussians

- Densities integrate to one:

$$\int_{-\infty}^{\infty} p(x; \mu, \sigma^2) dx = \int_{-\infty}^{\infty} \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right) dx = 1$$

- Mean:

$$\begin{aligned} \mathbb{E}_X[X] &= \int_{-\infty}^{\infty} xp(x; \mu, \sigma^2) dx \\ &= \int_{-\infty}^{\infty} x \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right) dx \\ &= \mu \end{aligned}$$

- Variance:

$$\begin{aligned} \mathbb{E}_X[(X - \mu)^2] &= \int_{-\infty}^{\infty} (x - \mu)^2 p(x; \mu, \sigma^2) dx \\ &= \int_{-\infty}^{\infty} (x - \mu)^2 \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right) dx \\ &= \sigma^2 \end{aligned}$$

$$X \sim \mathcal{N}(\mu, \sigma^2)$$


$$p(x; \mu, \sigma^2) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

Multivariate Gaussian

$$p(x; \mu, \Sigma) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} \exp \left(-\frac{1}{2} (x - \mu)^\top \Sigma^{-1} (x - \mu) \right)$$

$$\int \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} \exp \left(-\frac{1}{2} (x - \mu)^\top \Sigma^{-1} (x - \mu) \right) dx = 1$$

- Remember: For a matrix $A \in \mathbb{R}^{n \times n}$, $|A|$ denotes the determinant of A
- Remember: For a matrix $A \in \mathbb{R}^{n \times n}$, A^{-1} denotes the inverse of A
 - Rule: $A^{-1}A = I = AA^{-1}$
 - $I \in \mathbb{R}^{n \times n}$ is the identity matrix with all diagonal entries equal to one, and all off-diagonal entries equal to zero



Multivariate Gaussian

- Mean:

$$\mathbb{E}_X[X_i] = \int x_i p(x; \mu, \Sigma) dx = \mu_i$$


$$\mathbb{E}_X[X] = \int x p(x; \mu, \Sigma) dx = \mu \quad \text{(integral of vector = vector of integrals of each entry)}$$

- Covariance:


$$\mathbb{E}_X[(X_i - \mu_i)(X_j - \mu_j)] = \int (x_i - \mu_i)(x_j - \mu_j) p(x; \mu, \Sigma) dx = \Sigma_{ij}$$

$$\mathbb{E}_X[(X - \mu)(X - \mu)^\top] = \int [(X - \mu)(X - \mu)^\top] p(x; \mu, \Sigma) dx = \Sigma \quad \text{(integral of matrix = matrix of integrals of each entry)}$$


Multivariate Gaussian: examples



- $\mu = [1; 0]$
- $\Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$




- $\mu = [-.5; 0]$
- $\Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$




- $\mu = [-1; -1.5]$
- $\Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$

Multivariate Gaussian: examples




■ $\mu = [0; 0]$

■ $\Sigma = [1 \ 0; 0 \ 1]$



■ $\mu = [0; 0]$


■ $\Sigma = [.6 \ 0; 0 \ .6]$




■ $\mu = [0; 0]$

■ $\Sigma = [2 \ 0; 0 \ 2]$


Multivariate Gaussian: examples



- $\mu = [0; 0]$
- $\Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$




- $\mu = [0; 0]$
- $\Sigma = \begin{bmatrix} 1 & 0.5 \\ 0.5 & 1 \end{bmatrix}$




- $\mu = [0; 0]$
- $\Sigma = \begin{bmatrix} 1 & 0.8 \\ 0.8 & 1 \end{bmatrix}$


Multivariate Gaussian: examples





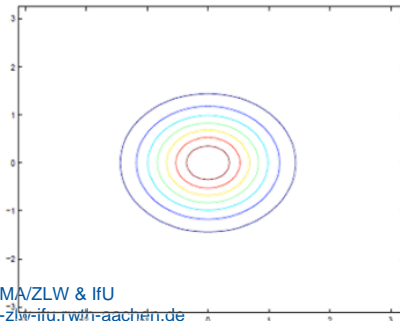
- $\mu = [0; 0]$
- $\Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$




- $\mu = [0; 0]$
- $\Sigma = \begin{bmatrix} 1 & 0.5 \\ 0.5 & 1 \end{bmatrix}$




- $\mu = [0; 0]$
- $\Sigma = \begin{bmatrix} 1 & 0.8 \\ 0.8 & 1 \end{bmatrix}$




Multivariate Gaussian: examples



- $\mu = [0; 0]$
- $\Sigma = \begin{bmatrix} 1 & -0.5 \\ -0.5 & 1 \end{bmatrix}$



- $\mu = [0; 0]$
- $\Sigma = \begin{bmatrix} 1 & -0.8 \\ -0.8 & 1 \end{bmatrix}$



- $\mu = [0; 0]$
- $\Sigma = \begin{bmatrix} 3 & 0.8 \\ 0.8 & 1 \end{bmatrix}$

Conditioning a multivariate Gaussian

- Consider a multi-variate Gaussian $\mathcal{N}(\mu, \Sigma) = \mathcal{N}\left(\begin{bmatrix} \mu_X \\ \mu_Y \end{bmatrix}, \begin{bmatrix} \Sigma_{XX} & \Sigma_{XY} \\ \Sigma_{YX} & \Sigma_{YY} \end{bmatrix}\right)$

- And the precision matrix $\Gamma = \Sigma^{-1} = \begin{bmatrix} \Sigma_{XX} & \Sigma_{XY} \\ \Sigma_{YX} & \Sigma_{YY} \end{bmatrix}^{-1} = \begin{bmatrix} \Gamma_{XX} & \Gamma_{XY} \\ \Gamma_{YX} & \Gamma_{YY} \end{bmatrix}$

- And the partitioned multi-variate Gaussian:

$$p\left(\begin{bmatrix} x \\ y \end{bmatrix}; \mu, \Sigma\right) = \frac{1}{(2\pi)^{(n/2)}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}\left(\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} \mu_X \\ \mu_Y \end{bmatrix}\right)^\top \begin{bmatrix} \Gamma_{XX} & \Gamma_{XY} \\ \Gamma_{YX} & \Gamma_{YY} \end{bmatrix} \left(\begin{bmatrix} x \\ y \end{bmatrix} - \begin{bmatrix} \mu_X \\ \mu_Y \end{bmatrix}\right)\right)$$

- We have

$$\begin{aligned} p(x|Y=y_0) &\propto p\left(\begin{bmatrix} x \\ y_0 \end{bmatrix}; \mu, \Sigma\right) \\ &\propto \exp\left(-\frac{1}{2}(x - \mu_X)^\top \Gamma_{XX}(x - \mu_X) - (x - \mu_X)^\top \Gamma_{XY}(y_0 - \mu_Y) - \frac{1}{2}(y_0 - \mu_Y)^\top \Gamma_{YY}(y_0 - \mu_Y)\right) \\ &\propto \exp\left(-\frac{1}{2}(x - \mu_X)^\top \Gamma_{XX}(x - \mu_X) - (x - \mu_X)^\top \Gamma_{XY}(y_0 - \mu_Y)\right) \\ &= \exp\left(-\frac{1}{2}(x - \mu_X)^\top \Gamma_{XX}(x - \mu_X) - (x - \mu_X)^\top \Gamma_{XX} \Gamma_{XX}^{-1} \Gamma_{XY}(y_0 - \mu_Y) - \frac{1}{2}(y_0 - \mu_Y)^\top \Gamma_{YX} \Gamma_{XX}^{-1} \Gamma_{XX} \Gamma_{XX}^{-1} \Gamma_{XY}(y_0 - \mu_Y) + \frac{1}{2}(y_0 - \mu_Y)^\top \Gamma_{YX} \Gamma_{XX}^{-1} \Gamma_{XX} \Gamma_{XX}^{-1} \Gamma_{XY}(y_0 - \mu_Y)\right) \\ &= \exp\left(-\frac{1}{2}(x - \mu_X + \Gamma_{XX}^{-1} \Gamma_{XY}(y_0 - \mu_Y))^\top \Gamma_{XX}(x - \mu_X + \Gamma_{XX}^{-1} \Gamma_{XY}(y_0 - \mu_Y))\right) \exp\left(\frac{1}{2}(y_0 - \mu_Y)^\top \Gamma_{YX} \Gamma_{XX}^{-1} \Gamma_{XX} \Gamma_{XX}^{-1} \Gamma_{XY}(y_0 - \mu_Y)\right) \\ &\propto \exp\left(-\frac{1}{2}(x - \mu_X + \Gamma_{XX}^{-1} \Gamma_{XY}(y_0 - \mu_Y))^\top \Gamma_{XX}(x - \mu_X + \Gamma_{XX}^{-1} \Gamma_{XY}(y_0 - \mu_Y))\right) \end{aligned}$$

Conditioning a multivariate Gaussian

- If

$$(X, Y) \sim \mathcal{N}(\mu, \Sigma) = \mathcal{N}\left(\begin{bmatrix} \mu_X \\ \mu_Y \end{bmatrix}, \begin{bmatrix} \Sigma_{XX} & \Sigma_{XY} \\ \Sigma_{YX} & \Sigma_{YY} \end{bmatrix}\right)$$

- Then:
$$\begin{aligned} X|Y = y_0 &\sim \mathcal{N}(\mu_X - \Gamma_{XX}^{-1}\Gamma_{XY}(y_0 - \mu_Y), \Gamma_{XX}) \\ &= \mathcal{N}(\mu_X + \Sigma_{XY}\Sigma_{YY}^{-1}(y_0 - \mu_Y), \Sigma_{XX} - \Sigma_{XY}\Sigma_{YY}^{-1}\Sigma_{YX}) \end{aligned}$$

- Mean moved according to correlation and variance on measurement
- Covariance $\Sigma_{XX|Y=y_0}$ does not depend on y_0

I. Gaussians

- I. Univariate Gaussian
- II. Multivariate Gaussian
- III. Conditioning (Bayes' rule)


II. Gaussian Mixture Model

III. Learning Local Dynamic Models

Disclaimer: lots of linear algebra now. In fact, pretty much all computations with Gaussians will be reduced to linear algebra!




Clustering

- Unsupervised Learning
- Detect patterns in unlabelled data
- Useful if you do not know what to look for
- Requires data, but no labels



Clustering

- Fundamental approach: Group **similar** instances
- Example: 2D pattern



- What does **similarity** mean?

What does similarity mean?

- Similarity is hard to define
 - But we know it if we see it
- The true meaning of similarity is a philosophical question. We will therefore choose a more pragmatic approach: we think of **distances** (rather than similarities) between vectors or correlations between random variables



Hard Clustering with K-Means

K-Means Algorithm




1. Choose K random points as cluster centers
2. Assign datapoint to the nearest cluster center
3. Adjust cluster center so that it gets the mean value of the associated points

Problem: correct assignment is hard!

- Sometimes distances can be **deceiving!**
- Cluster may overlap!

What about probabilistic clustering?



The General Gaussian Mixture Model Assumption


- “The probabilistic version of K-Means”
- Each cluster has not only mean μ_i but also associated covariance matrix Σ_i
- GMM is a linear combination of K Gaussians given by

$$p(\mathbf{x}) = \sum_{k=1}^K p(\mathbf{x}|k)P(k)$$

where $P(k)$ is mixture weight subject to constraints

$$0 \leq P(k) \leq 1 \quad \text{and} \quad \sum_{k=1}^K P(k) = 1$$


and $p(x|k)$ is height of k 'th Gaussian at datapoint x




The General Gaussian Mixture Model Assumption

- Gaussian Mixture Model
 - $P(Y)$ is multinomial
 - $p(x|k)$ is a multivariate Gaussian Distribution

$$P(X = \mathbf{x}_j | Y = i) = \frac{1}{(2\pi)^{m/2} \|\Sigma_i\|^{1/2}} \exp\left[-\frac{1}{2}(\mathbf{x}_j - \mu_i)^T \Sigma_i^{-1} (\mathbf{x}_j - \mu_i)\right]$$



Single Gaussian




Mixture of two Gaussians

Combine simple models into a complex model




$$p(\mathbf{x}) = \sum_{k=1}^K \pi_k \underbrace{\mathcal{N}(\mathbf{x} | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}_{\text{Component}}$$

Mixing coefficient

$$\forall k : \pi_k \geq 0 \quad \sum_{k=1}^K \pi_k = 1$$




Combine simple models into a complex model: Example



Soft-Clustering with Expectation Maximization and GMM

Expectation-Maximization Algorithm


1. Choose K random mean and covariances as clusters
2. E-Step: Calculate, for each point, the probabilities of it belonging to each of the clusters
3. M-Step: recalculate mean and covariance of each cluster, using the probability of belonging to each cluster



$$p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t) = \mathcal{N}(f(\mathbf{x}_t, \mathbf{u}_t), \Sigma)$$

$$f(\mathbf{x}_t, \mathbf{u}_t) \approx \mathbf{A}_t \mathbf{x}_t + \mathbf{B}_t \mathbf{u}_t$$

$$\mathbf{A}_t = \frac{\partial f}{\partial \mathbf{x}_t} \quad \mathbf{B}_t = \frac{\partial f}{\partial \mathbf{u}_t}$$



Linearized local dynamics

Goal: get the system dynamics $p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t)$ for each timestep t

Data: samples generated by the previous controller $\hat{p}_i(\mathbf{u}_t|\mathbf{x}_t) \rightarrow \{(\mathbf{x}_t, \mathbf{u}_t, \mathbf{x}_{t+1})_i\}$

Linear Gaussian Dynamics are defined as

$$p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t) = \mathcal{N}(f_{xt}\mathbf{x}_t + f_{ut}\mathbf{u}_t + f_{ct}), \mathbf{F}_t)$$

How can we determine linear
Gaussian dynamics from few
samples?

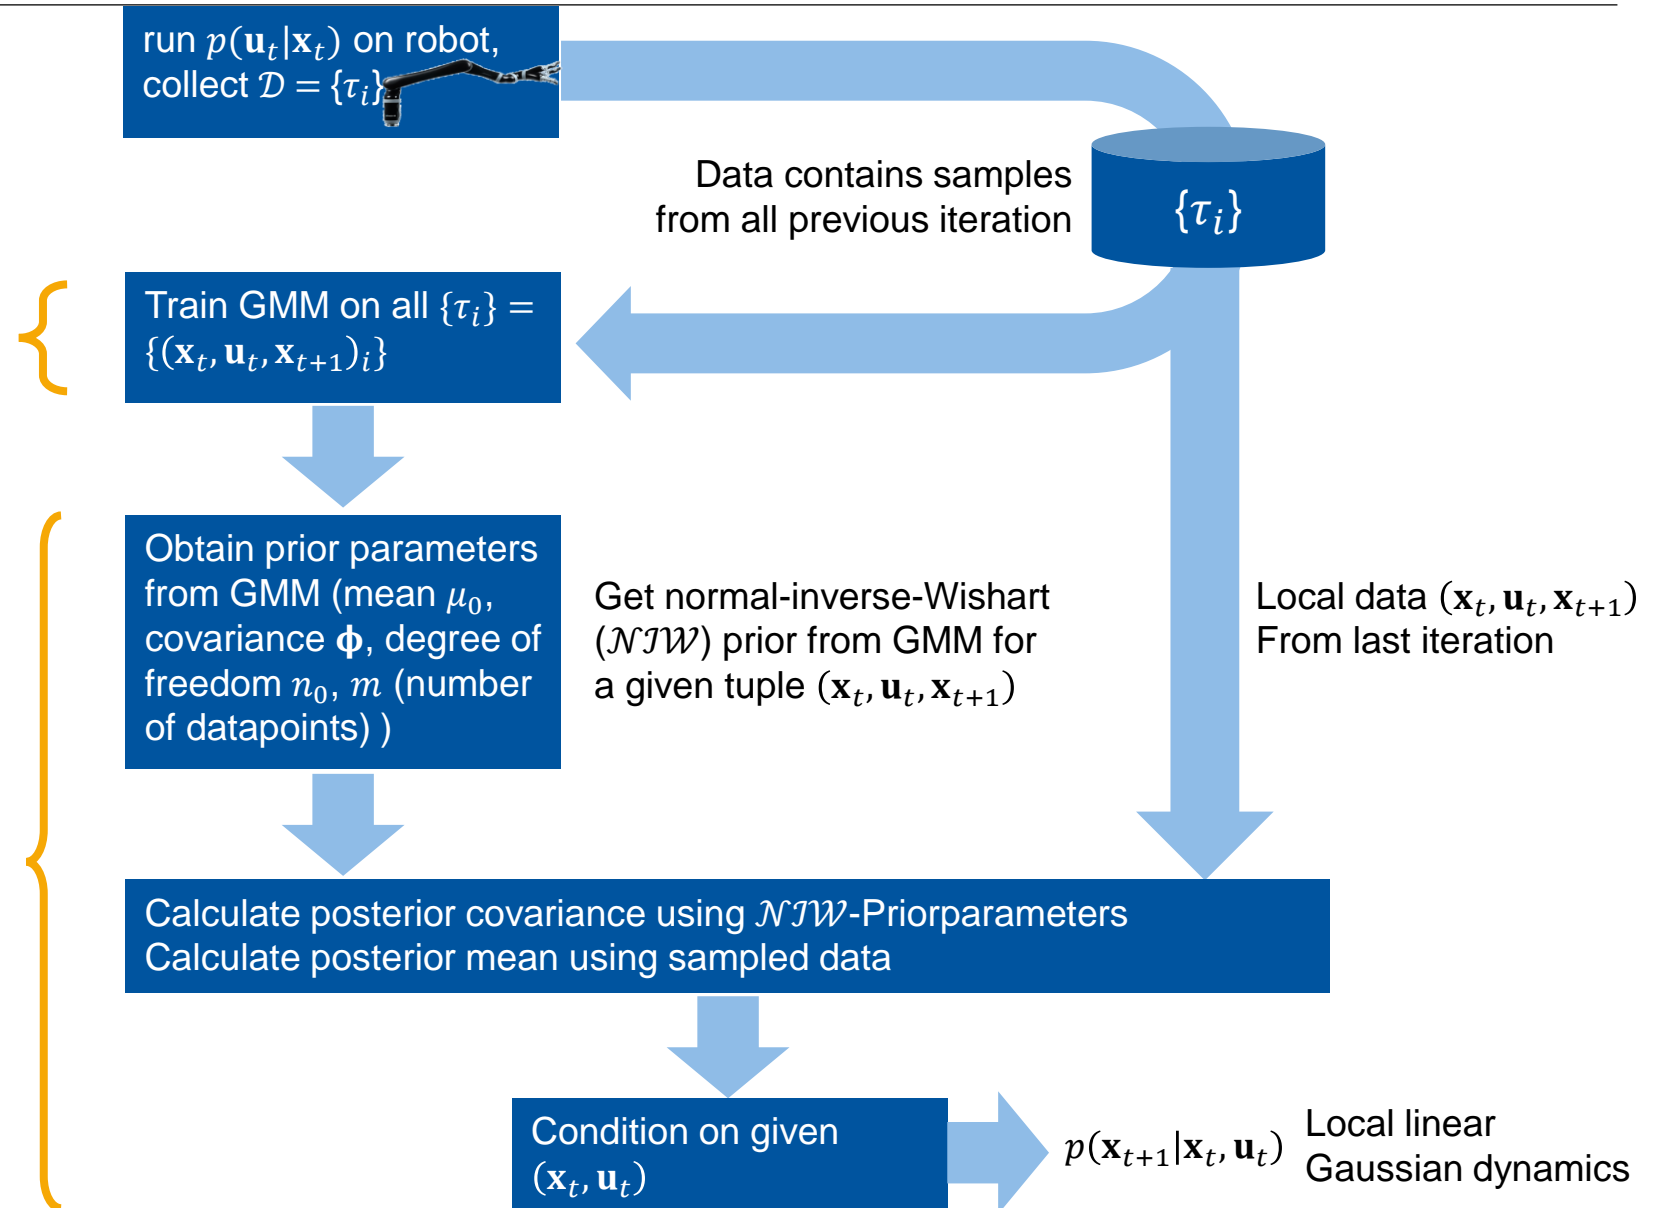
Learning Local and Global Models

Train GMM: **Global Dynamic Model**

- Uses data from nearby timesteps
- Uses data from prior iterations

Linearize: **Local Dynamic Model**

- Uses prior from local dynamic model
- Uses data from last iteration at timestep t
- Condition on given $(\mathbf{x}_t, \mathbf{u}_t)$



Its your turn!




Visit the website and implement it!



Introduction to the tasks

Tasks for today and tomorrow

- Task 1:
 - Implement an LQR Backward and Forward pass
 - Try to understand it!
 - Test it with our test method
- Task 2:
 - Implement linearization of the dynamic model
 - Try to understand it!
 - Test it with our test-method
 - Test it on the Box2D Scenario
- Task 3:
 - Test it with Kinova Jaco 2 in simulation
 - Adjust cost function



Task 2 – Installation procedure

Download source code (do it in your home directory: `cd ~`):

```
git clone https://github.com/philippente/task2\_dynamics.git
```

Edit `.bashrc` to set environment variables:

```
gedit ~/.bashrc
```

At the end of file, the lines should look like this:

```
source /opt/ros/indigo/setup.bash
```

```
source /home/<USERNAME>/catkin_ws/devel/setup.bash
```

```
export
```

```
ROS_PACKAGE_PATH=$ROS_PACKAGE_PATH:/opt/ros/indigo/share:/opt/ros/indigo/stacks:/home/<USERNAME>/task2_dynamics:/home/<USERNAME>/task2_dynamics/src/gps_agent_pkg
```

Check if the blue part of the source folder and `ROS_PACKAGE_PATH` is correct!

Then save it and close it. Source the `.bashrc` (load the environment variables):

```
source ~/.bashrc
```

Now, compile some stuff:

```
cd task2_dynamics
```

```
catkin_make
```


Task 2 – Installation procedure

- Open PyCharm
- Import the folder *task2_dynamics* as a new project
- Open within PyCharm: `python/gps/algorithm/dynamics/dynamics_lr_prior.py`
- **Task: Implement dynamics learning! Look at the website for advices: <https://philippente.github.io/irobotics.html>**
- You can test your implementation with a little test program
 - using a terminal, open the directory *task2_dynamics*
 - Start the program with: `python python/gps/dynamics_test.py`
 - Was it successful?
- If it was successful, lets look at the Box2D scenario!
 - using a terminal, open the directory *task2_dynamics*
 - Start it with `python python/gps/gps_main.py box2d_arm_example`
 - How is the performance?

