



Source: <https://blenderartists.org/forum/showthread.php?316399-Industrial-Robots>

Intelligent Robotics

Summary – Intelligent Motion Planning

Philipp Ennen, M.Sc.

Robots as an Example for Intelligent Machines

The central questions of robotics

Three main question:

Where am I?



- Easy in Industrial Robotics
- Hard in Mobile Robotics


Where should I go?



- Often hand-engineered
- Or a Result of Task Planning

This weeks focus!

How do I go there?

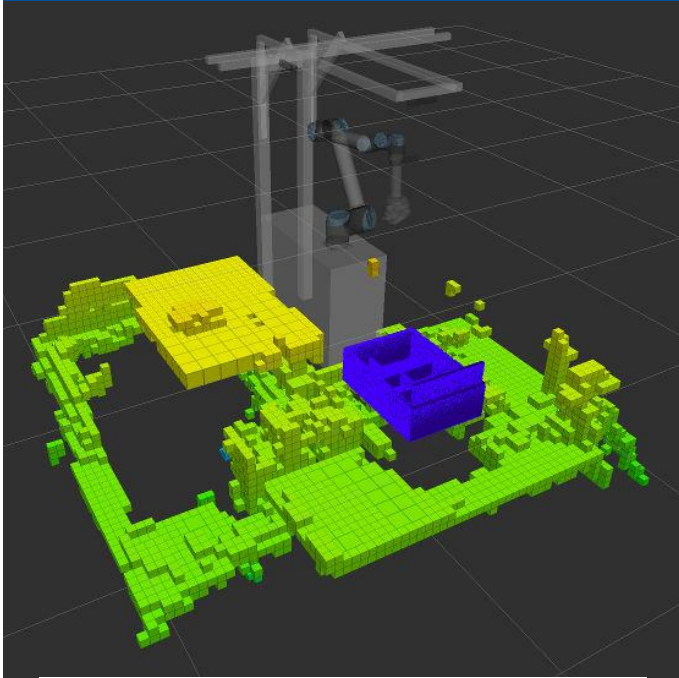


- Comparable Easy in Mobile Robotics
- Hard in Industrial Robotics


Robots as an Example for Intelligent Machines

How do I (the robot) go there?

Model of the Robot and Environment



Sensing




This weeks topic!

Motion Planning

Requires Goalstate:

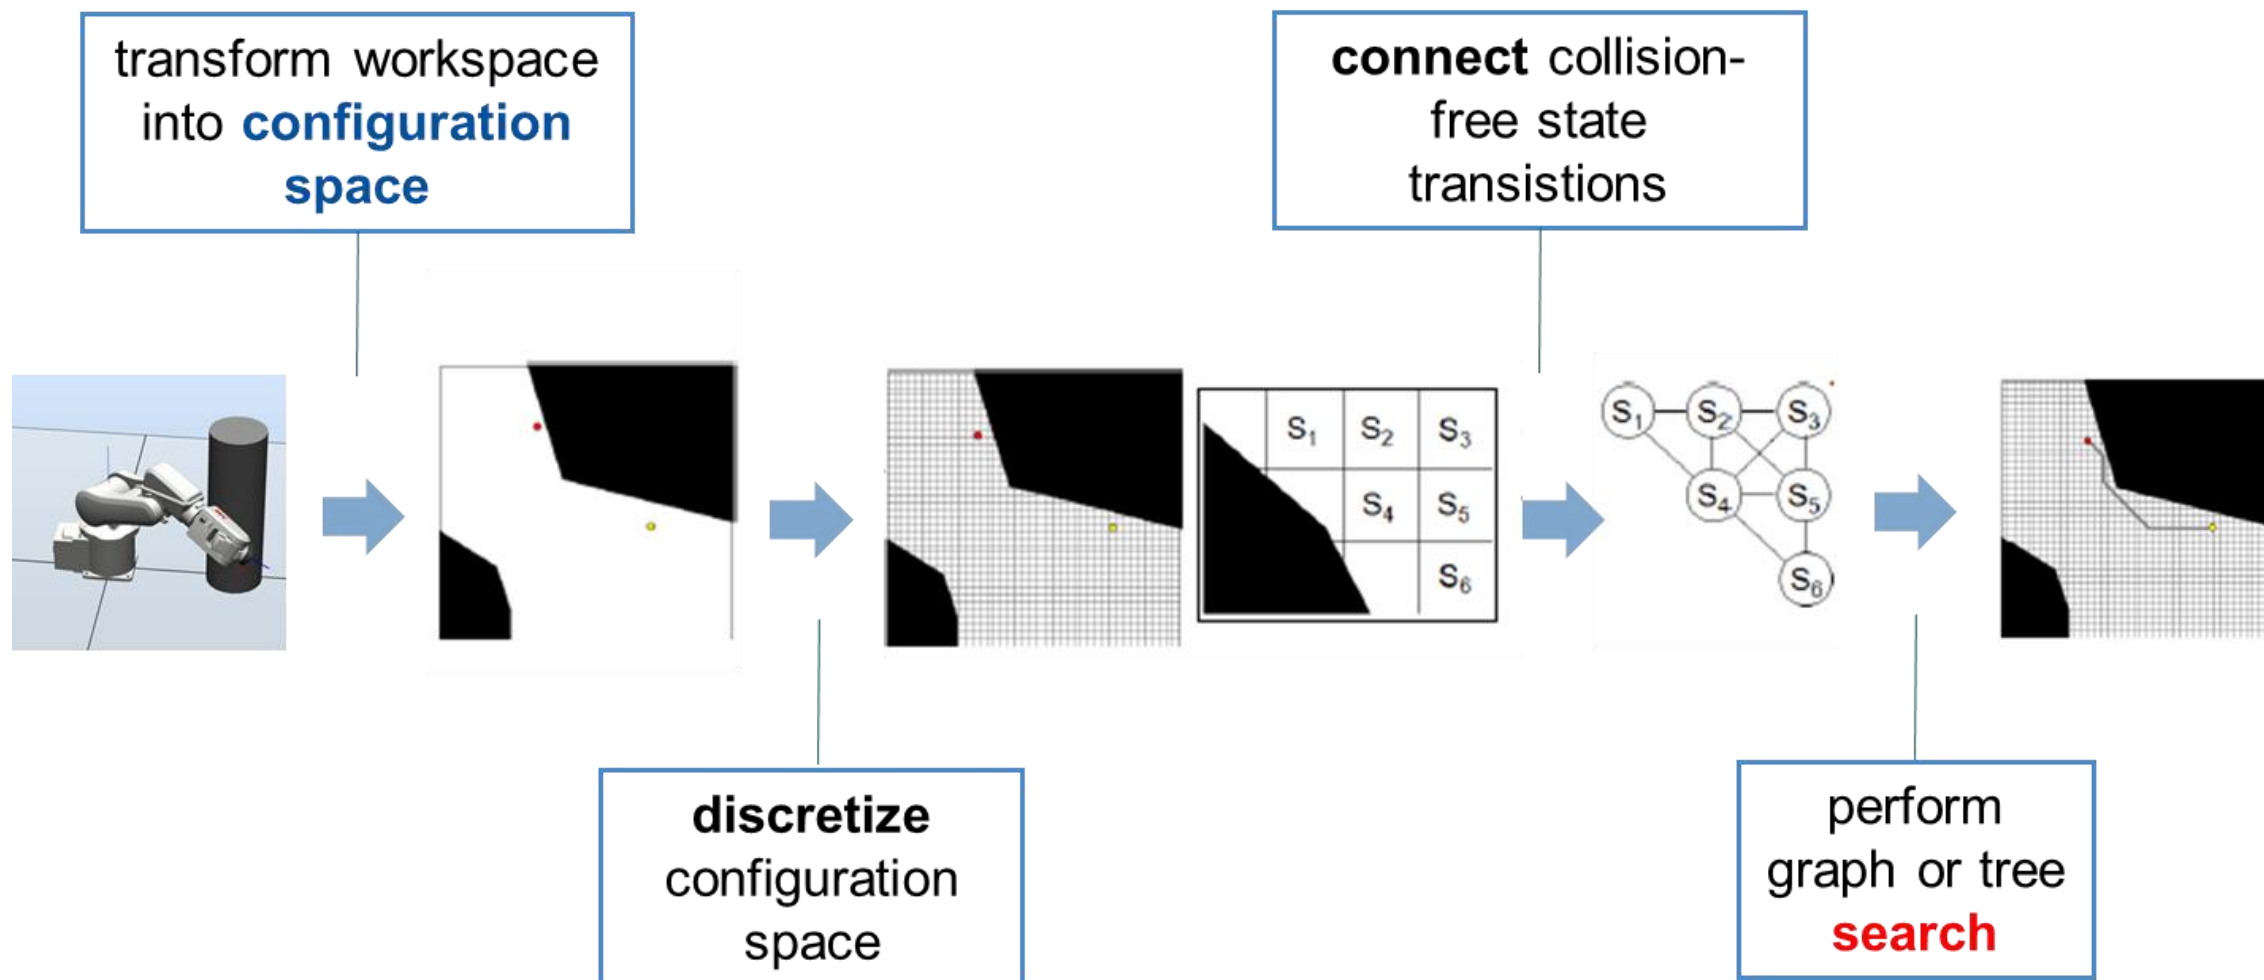
- i.e. hand-engineered
- i.e. via a cost function

Motion Control and Execution



Motion Planning: Configuration Space

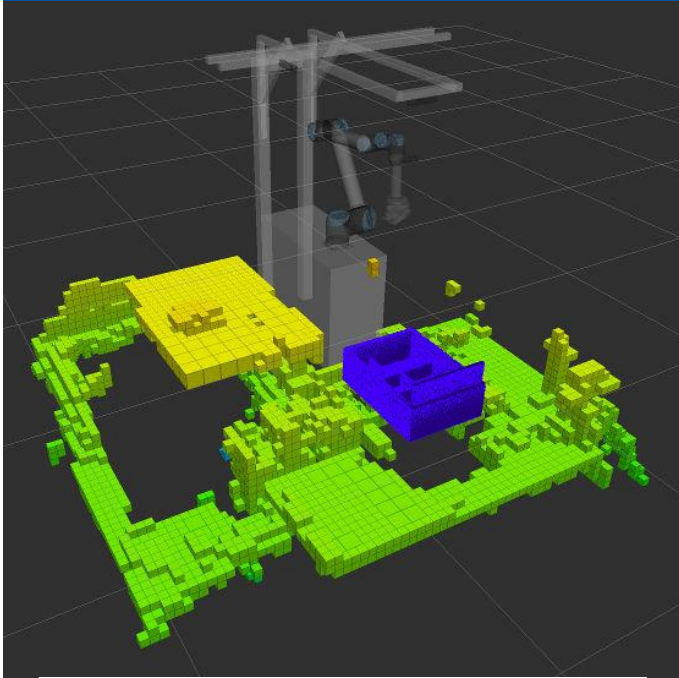
Motion Planning Pipeline within \mathcal{C}_{space}



Robots as an Example for Intelligent Machines


How do I (the robot) go there?

Model of the Robot and Environment



$$\begin{bmatrix} \ddot{x} \\ \ddot{y} \\ \ddot{\phi} \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \frac{-(I + ml^2)b}{I(M+m) + Mmi^2} & \frac{m^2 gl^2}{I(M+m) + Mmi^2} \\ 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ \phi \end{bmatrix} + \begin{bmatrix} 0 \\ \frac{I + ml^2}{I(M+m) + Mmi^2} \\ \frac{mgl(M+m)}{I(M+m) + Mmi^2} \end{bmatrix} u$$

Sensing




What if the model of the robot and environment is hard to describe (or unknown)?

Think about flexible objects!



Motion Control and Execution

Think about contact-situations!




Requires Goalstate:

- i.e. hand-engineered
- i.e. via a cost function

Robots as an Example for Intelligent Machines

What if the model of the robot and environment is hard to describe (or unknown)?

- Use an Reinforcement Learning Agent!




Robots as an Example for Intelligent Machines

What if the model of the robot and environment is hard to describe (or unknown)?

This weeks topic!

Model-based RL:

- Learn to predict next state: $P(s'|s, a)$
- Learn to predict immediate reward $P(r'|s, a)$



Model-free RL:

- Learn to predict value: $V(s)$ or $Q(s, a)$

Linear Quadratic Regulator

- Special case: Systems with
 - **Linear dynamics**
 - **Quadratic costs**
- LQR provides an exact solution

$$\min_{\mathbf{u}_1, \dots, \mathbf{u}_T} c(\mathbf{x}_1, \mathbf{u}_1) + c(f(\mathbf{x}_1, \mathbf{u}_1), \mathbf{u}_2) + \dots + c(f(f(\dots)), \mathbf{u}_T)$$

$$f(\mathbf{x}_t, \mathbf{u}_t) = \mathbf{F}_t \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix} + \mathbf{f}_t$$

linear

$$c(\mathbf{x}_t, \mathbf{u}_t) = \frac{1}{2} \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{C}_t \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix} + \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{c}_t$$

quadratic

Linear Quadratic Regulator

Pseudocode Algorithm

Backward recursion
(Get linear equations for \mathbf{u})

for $t = T$ to 1:

$$\mathbf{Q}_t = \mathbf{C}_t + \mathbf{F}_t^T \mathbf{V}_{t+1} \mathbf{F}_t$$

$$\mathbf{q}_t = \mathbf{c}_t + \mathbf{F}_t^T \mathbf{V}_{t+1} \mathbf{f}_t + \mathbf{F}_t^T \mathbf{v}_{t+1}$$

$$Q(\mathbf{x}_t, \mathbf{u}_t) = \text{const} + \frac{1}{2} \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{Q}_t \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix} + \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{q}_t$$

$$\mathbf{u}_t \leftarrow \arg \min_{\mathbf{u}_t} Q(\mathbf{x}_t, \mathbf{u}_t) = \mathbf{K}_t \mathbf{x}_t + \mathbf{k}_t$$


$$\mathbf{K}_t = -\mathbf{Q}_{\mathbf{u}_t, \mathbf{u}_t}^{-1} \mathbf{Q}_{\mathbf{u}_t, \mathbf{x}_t}$$

$$\mathbf{k}_t = -\mathbf{Q}_{\mathbf{u}_t, \mathbf{u}_t}^{-1} \mathbf{q}_{\mathbf{u}_t}$$

$$\mathbf{V}_t = \mathbf{Q}_{\mathbf{x}_t, \mathbf{x}_t} + \mathbf{Q}_{\mathbf{x}_t, \mathbf{u}_t} \mathbf{K}_t + \mathbf{K}_t^T \mathbf{Q}_{\mathbf{u}_t, \mathbf{x}_t} \mathbf{x}_t + \mathbf{K}_t^T \mathbf{Q}_{\mathbf{u}_t, \mathbf{u}_t} \mathbf{K}_t$$

$$\mathbf{v}_t = \mathbf{K}_t^T \mathbf{Q}_{\mathbf{u}_t, \mathbf{u}_t} \mathbf{k}_t + \mathbf{Q}_{\mathbf{x}_t, \mathbf{u}_t} \mathbf{k}_t + \mathbf{q}_{\mathbf{x}_t} + \mathbf{K}_t^T \mathbf{q}_{\mathbf{u}_t}$$

$$V(\mathbf{x}_t) = \text{const} + \frac{1}{2} \mathbf{x}_t^T \mathbf{V}_t \mathbf{x}_t + \mathbf{x}_t^T \mathbf{v}_t$$



Forward recursion
(Use known initial state to get
values for \mathbf{u})

for $t = 1$ to T :

$$\mathbf{u}_t = \mathbf{K}_t \mathbf{x}_t + \mathbf{k}_t$$

$$\mathbf{x}_{t+1} = f(\mathbf{x}_t, \mathbf{u}_t)$$

Why do we want to learn the dynamics?

- If $\mathbf{x}_{t+1} = f(\mathbf{x}_t, \mathbf{u}_t)$ is known, we can do trajectory optimization
 - In the stochastic case $p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t)$



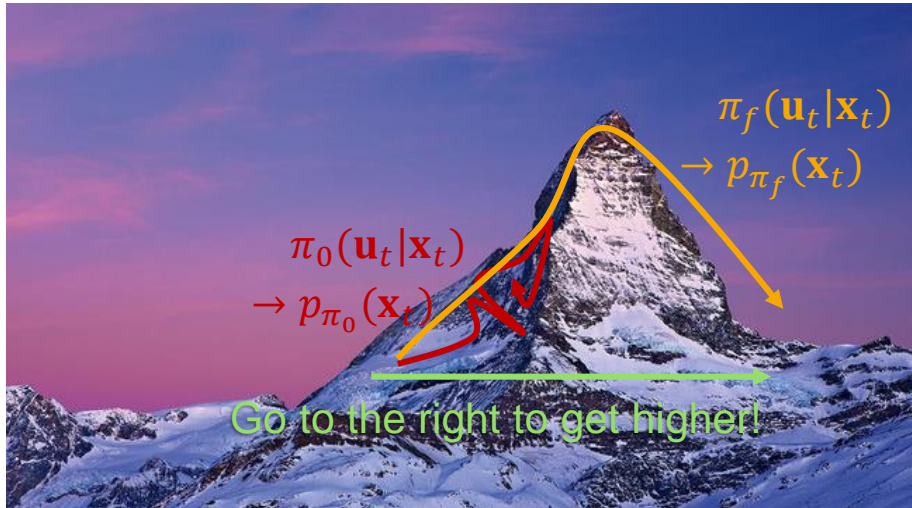
Learn $f(\mathbf{x}_t, \mathbf{u}_t)$ with subsequent backpropagation (i.e. iLQR)

Modelbased Reinforcement Learning Version 0.5

1. Execute initial policy $\pi_0(\mathbf{u}_t|\mathbf{x}_t)$ (i.e. a random policy) and collect data $\mathcal{D} = \{(\mathbf{x}, \mathbf{u}, \mathbf{x}')_i\}$
2. Learn dynamics $f(\mathbf{x}, \mathbf{u})$ that minimizes $\sum_i \|f(\mathbf{x}_i, \mathbf{u}_i) - \mathbf{x}'_i\|^2$
3. Backpropagate $f(\mathbf{x}, \mathbf{u})$ and calculate sequence of actions (i.e. iLQR)

Does Version 0.5 work?

(in general) **NO!**



1. Execute initial policy $\pi_0(\mathbf{u}_t|\mathbf{x}_t)$ (i.e. a random policy) and collect data $\mathcal{D} = \{(\mathbf{x}, \mathbf{u}, \mathbf{x}')_i\}$
2. Learn dynamics $f(\mathbf{x}, \mathbf{u})$ that minimizes $\sum_i \|f(\mathbf{x}_i, \mathbf{u}_i) - \mathbf{x}'_i\|^2$
3. Backpropagate $f(\mathbf{x}, \mathbf{u})$ and calculate sequence of actions (i.e. iLQR) $\rightarrow \pi_f(\mathbf{u}_t|\mathbf{x}_t)$

$$p_{\pi_0}(\mathbf{x}_t) \neq p_{\pi_f}(\mathbf{x}_t)$$

(Distribution Mismatch Problem)



Distribution Mismatch Problem increases if expressive classes of models are used (i.e. neural networks)

Can we do better?

Can we make $p_{\pi_0}(\mathbf{x}_t) = p_{\pi_f}(\mathbf{x}_t)$?



Need to collect data from $p_{\pi_f}(\mathbf{x}_t)$!


Modellbasiertes Reinforcement Learning Version 1.0

1. Execute initial policy $\pi_0(\mathbf{u}_t|\mathbf{x}_t)$ (i.e. a random policy) and collect data $\mathcal{D} = \{(\mathbf{x}, \mathbf{u}, \mathbf{x}')_i\}$
2. Learn dynamics $f(\mathbf{x}, \mathbf{u})$ that minimizes $\sum_i \|f(\mathbf{x}_i, \mathbf{u}_i) - \mathbf{x}'_i\|^2$
3. Backpropagate $f(\mathbf{x}, \mathbf{u})$ and calculate sequence of actions (i.e. iLQR)
4. Execute those actions and add the resulting data $\{(\mathbf{x}, \mathbf{u}, \mathbf{x}')_i\}$ to \mathcal{D}

$$p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t) = \mathcal{N}(f(\mathbf{x}_t, \mathbf{u}_t), \Sigma)$$

$$f(\mathbf{x}_t, \mathbf{u}_t) \approx \mathbf{A}_t \mathbf{x}_t + \mathbf{B}_t \mathbf{u}_t$$

$$\mathbf{A}_t = \frac{\partial f}{\partial \mathbf{x}_t} \quad \mathbf{B}_t = \frac{\partial f}{\partial \mathbf{u}_t}$$



Linearized local dynamics

Goal: get the system dynamics $p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t)$ for each timestep t

Data: samples generated by the previous controller $\hat{p}_i(\mathbf{u}_t|\mathbf{x}_t) \rightarrow \{(\mathbf{x}_t, \mathbf{u}_t, \mathbf{x}_{t+1})_i\}$


Linear Gaussian Dynamics are defined as

$$p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t) = \mathcal{N}(f_{xt}\mathbf{x}_t + f_{ut}\mathbf{u}_t + f_{ct}, \mathbf{F}_t)$$

How can we determine linear
Gaussian dynamics from few
samples?

What kind of models can we use?

Gaussian process



GP with input (\mathbf{x}, \mathbf{u}) and output \mathbf{x}'

Pro: very data-efficient

Con: not great with non-smooth dynamics

Con: very slow when dataset is big

Neural Network




image: Punjani & Abbeel '14


Input is (\mathbf{x}, \mathbf{u}) , output is \mathbf{x}'

Pro: very expressive, can use lots of data

Con: not so great in low data regimes

This week's focus!

Gaussian Mixture Model



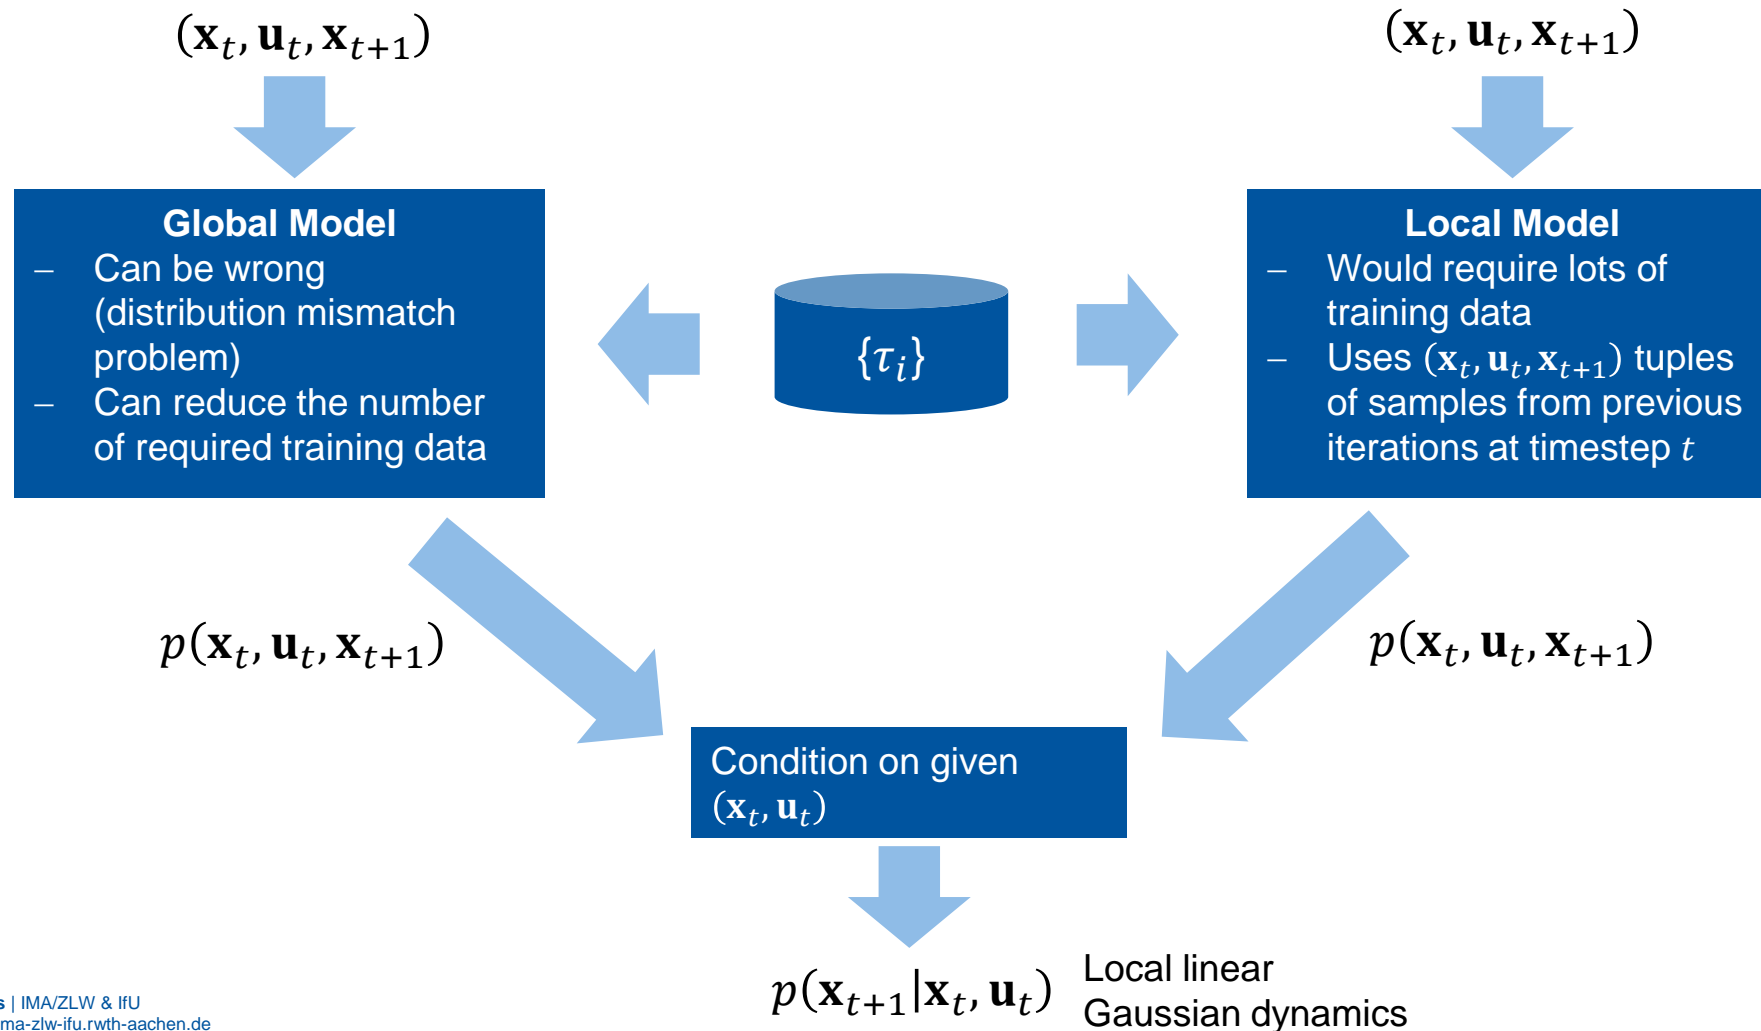
GMM over $(\mathbf{x}, \mathbf{u}, \mathbf{x}')$ tuples

Train on $(\mathbf{x}, \mathbf{u}, \mathbf{x}')$, condition to get $p(\mathbf{x}'|\mathbf{x}, \mathbf{u})$

For i 'th mixture element, $p_i(\mathbf{x}, \mathbf{u})$ gives region where the mode $p_i(\mathbf{x}'|\mathbf{x}, \mathbf{u})$ holds

Pro: very expressive, if the dynamics can be assumed as piecewise linear

Combining global and local models



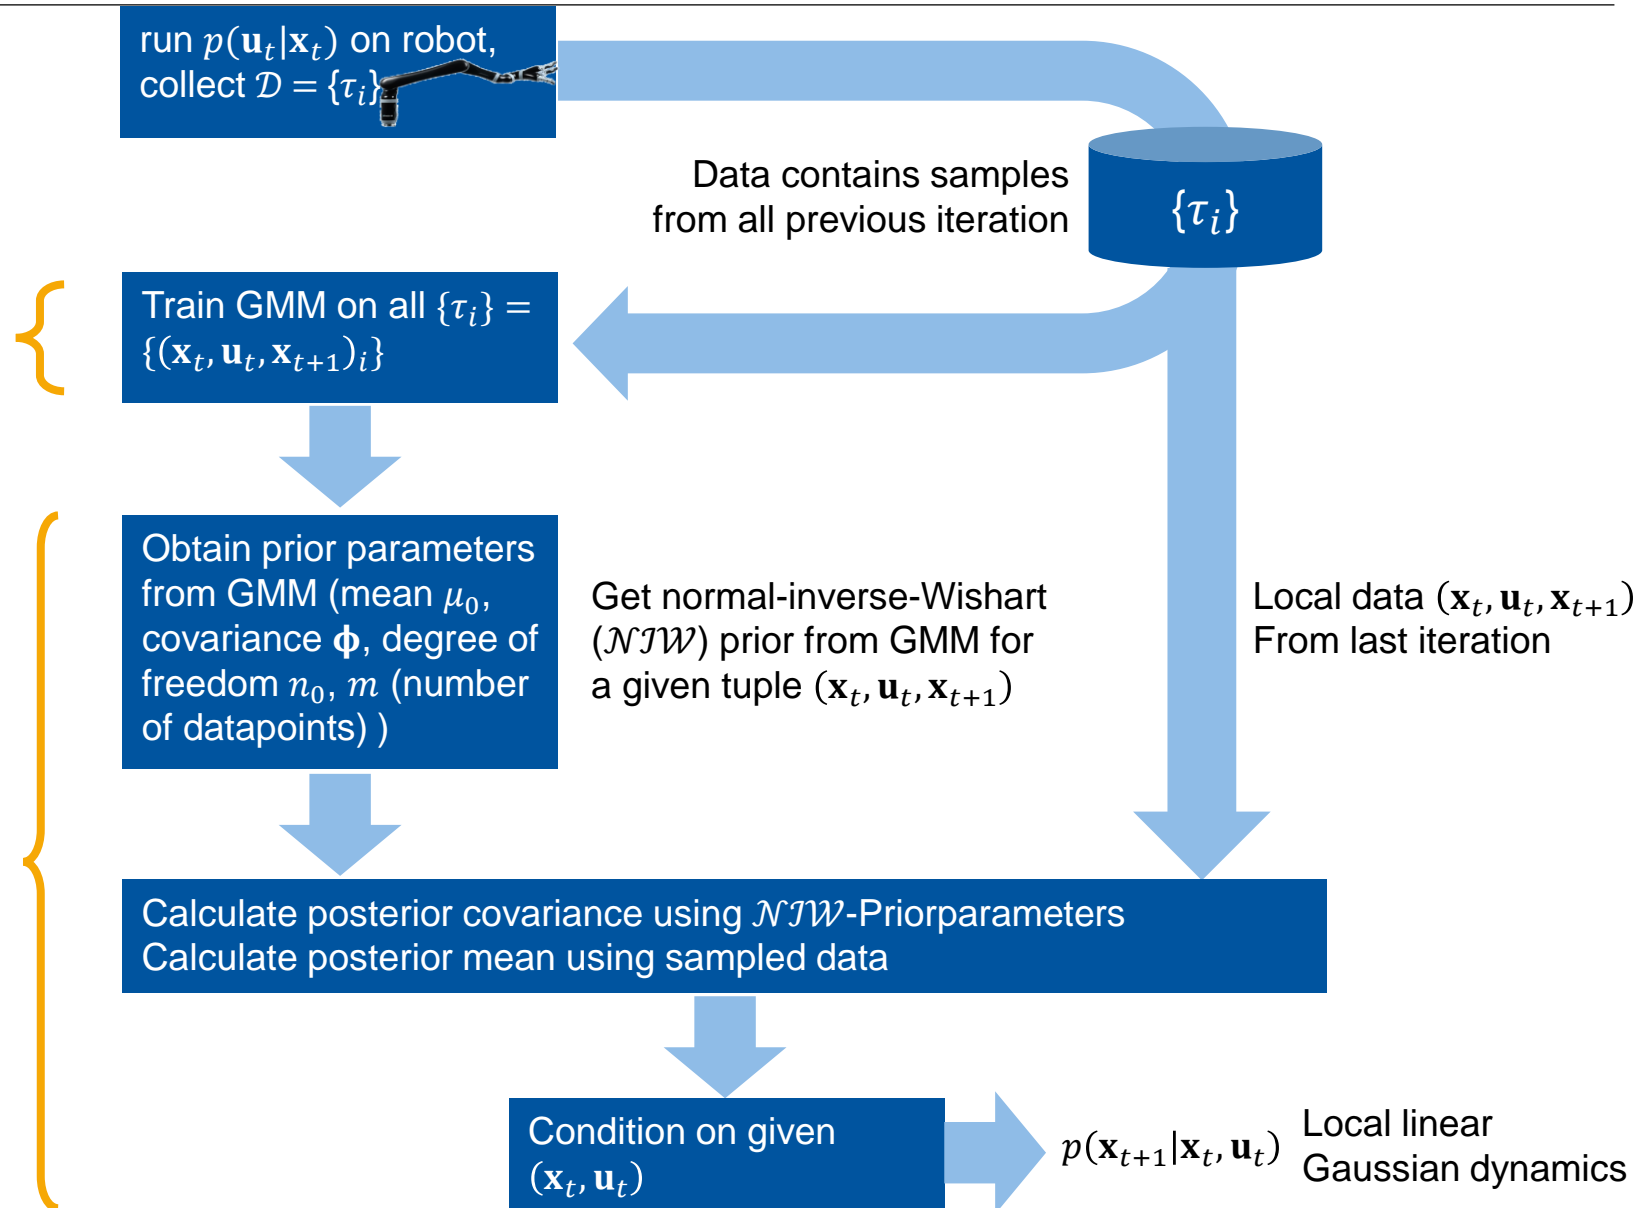
Learning Local and Global Models

Train GMM: **Global Dynamic Model**

- Uses data from nearby timesteps
- Uses data from prior iterations

Linearize: **Local Dynamic Model**

- Uses prior from local dynamic model
- Uses data from last iteration at timestep t
- Condition on given $(\mathbf{x}_t, \mathbf{u}_t)$



**Thank you
for your attention!**