

Source: https://ttechcrunch2011.files.wordpress.com/2016/02/shutterstock_147776027.jpg?w=1279&h=727&crop=1

Robotics for Future Industrial Applications

Modelbased Reinforcement Learning

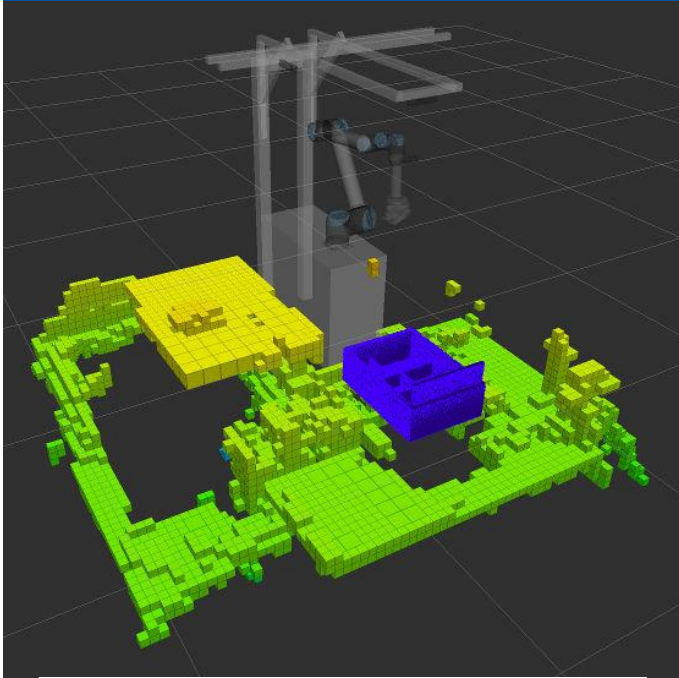
Philipp Ennen, M.Sc.




Introduction

How do I (the robot) go there?

Model of the Robot and Environment



Sensing

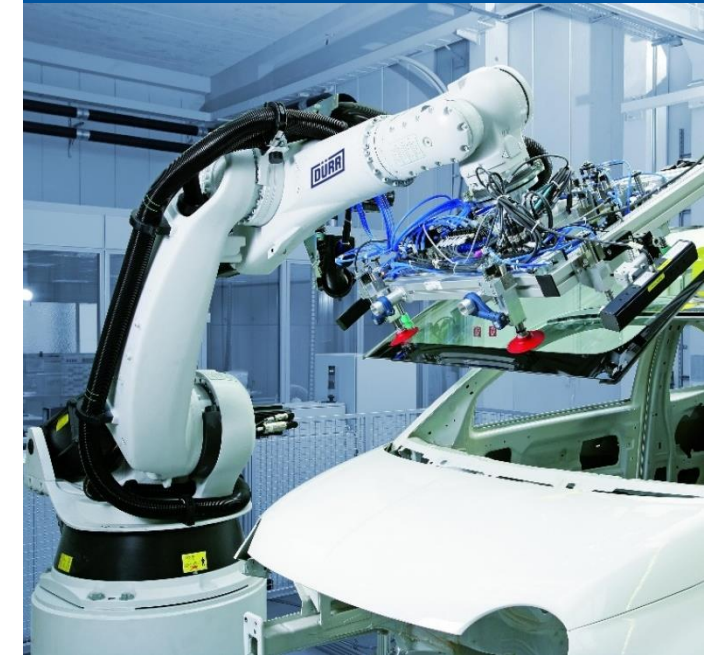


Motion Planning with iLQR

Requires Goalstate:

- i.e. hand-engineered
- i.e. via a cost function

Motion Control and Execution

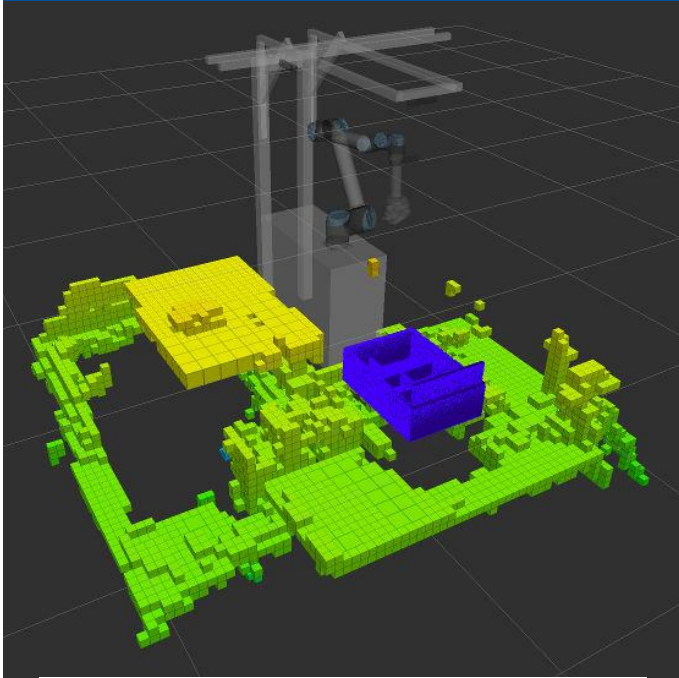


$$\begin{bmatrix} \ddot{x} \\ \ddot{\phi} \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & \frac{-(I + ml^2)b}{I(M+m) + Mmi^2} & \frac{m^2 gl^2}{I(M+m) + Mmi^2} & 0 \\ 0 & 0 & 0 & 1 \\ 0 & \frac{-mlb}{I(M+m) + Mmi^2} & \frac{mgl(M+m)}{I(M+m) + Mmi^2} & 0 \end{bmatrix} \begin{bmatrix} x \\ \dot{x} \\ \phi \\ \dot{\phi} \end{bmatrix} + \begin{bmatrix} 0 \\ \frac{I + ml^2}{I(M+m) + Mmi^2} \\ 0 \\ \frac{ml}{I(M+m) + Mmi^2} \end{bmatrix} u$$

Introduction

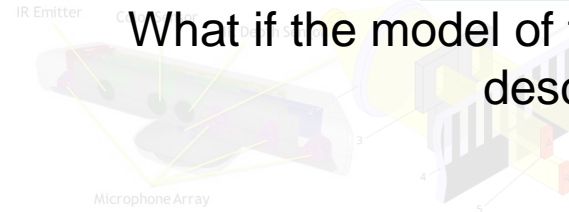
How do I (the robot) go there?

Model of the Robot and Environment



$$\begin{bmatrix} \ddot{x} \\ \ddot{\phi} \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & \frac{-(I + ml^2)b}{I(M+m) + Mmi^2} & \frac{m^2 gl^2}{I(M+m) + Mmi^2} & 0 \\ 0 & 0 & 0 & 1 \\ 0 & \frac{-mlb}{I(M+m) + Mmi^2} & \frac{mgl(M+m)}{I(M+m) + Mmi^2} & 0 \end{bmatrix} \begin{bmatrix} x \\ \dot{x} \\ \phi \\ \dot{\phi} \end{bmatrix} + \begin{bmatrix} 0 \\ \frac{I + ml^2}{I(M+m) + Mmi^2} \\ 0 \\ \frac{ml}{I(M+m) + Mmi^2} \end{bmatrix} u$$

Sensing



What if the model of the robot and environment is hard to describe (or unknown)?


Think about flexible objects!



Requires Goalstate:

- i.e. hand-engineered
- i.e. via a cost function

Motion Control and Execution




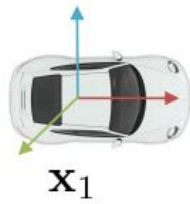
Think about contact-situations!

Dynamics of unstructured environments



Real-World Dynamics are Complex!

Often dynamic models exist



Dynamic models usually do not exist



Dynamics of contacts



Dynamics of flexible objects



Dynamics of unstructured environments

Trajectory Optimization

- Trajectory Optimization: Calculates **optimal sequence of actions** using **cost-function** and **dynamics**

$$\min_{\mathbf{u}_1, \dots, \mathbf{u}_T} \sum_{t=1}^T c(\mathbf{x}_t, \mathbf{u}_t) \quad \text{s.t.} \quad \mathbf{x}_t = f(\mathbf{x}_{t-1}, \mathbf{u}_{t-1})$$

- Today: How can optimal action sequences be calculated if a dynamic model does not exist?
- Today we learn an algorithm based on
 - Learning a global dynamic model (“model-based reinforcement learning”)
 - Learning a local dynamic model


Robots as an Example for Intelligent Machines

What if the model of the robot and environment is hard to describe (or unknown)?

This weeks topic!

Model-based RL:

- Learn to predict next state (using a dynamic model): $P(s' | s, a)$
- ~~Learn to predict immediate reward~~
 $P(r^+ | s, a)$ (we assume to have this information)



Model-free RL:

- Learn to predict value: $V(s)$ or $Q(s, a)$

I. Modelbased Reinforcement Learning

- I. Learning of dynamic models
- II. Learning of dynamic models and policies

II. Representing a dynamic model

III. Global and local dynamic model

IV. Learning with local dynamic models with „Trust Regions“

Why do we want to learn the dynamics?

$$\min_{\mathbf{u}_1, \dots, \mathbf{u}_T} \sum_{t=1}^T c(\mathbf{x}_t, \mathbf{u}_t) \quad \text{s.t.} \quad \mathbf{x}_t = f(\mathbf{x}_{t-1}, \mathbf{u}_{t-1})$$



$$\min_{\mathbf{u}_1, \dots, \mathbf{u}_T} c(\mathbf{x}_1, \mathbf{u}_1) + c(f(\mathbf{x}_1, \mathbf{u}_1), \mathbf{u}_2) + \dots + c(f(f(\dots) \dots), \mathbf{u}_T)$$

Usual procedure: Differentiate via Backpropagation and optimize (i.e. iLQR)

Requires: $\frac{\partial f}{\partial \mathbf{x}_t}, \frac{\partial f}{\partial \mathbf{u}_t}, \frac{\partial c}{\partial \mathbf{x}_t}, \frac{\partial c}{\partial \mathbf{u}_t}$

Why do we want to learn the dynamics?

- If $\mathbf{x}_{t+1} = f(\mathbf{x}_t, \mathbf{u}_t)$ is known, we can do trajectory optimization
 - In the stochastic case $p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t)$



Learn $f(\mathbf{x}_t, \mathbf{u}_t)$ with subsequent backpropagation (i.e. iLQR)

Modelbased Reinforcement Learning Version 0.5

1. Execute initial policy $\pi_0(\mathbf{u}_t|\mathbf{x}_t)$ (i.e. a random policy) and collect data $\mathcal{D} = \{(\mathbf{x}, \mathbf{u}, \mathbf{x}')_i\}$
2. Learn dynamics $f(\mathbf{x}, \mathbf{u})$ that minimizes $\sum_i \|f(\mathbf{x}_i, \mathbf{u}_i) - \mathbf{x}'_i\|^2$
3. Backpropagate $f(\mathbf{x}, \mathbf{u})$ and calculate sequence of actions (i.e. iLQR)


Does Version 0.5 work?

(often) **YES!**

- Traditional system identification uses this method (control theory)
- Initial policy must be chosen with caution
- Version 0.5 is very effective
 - If a representation of the dynamics based on physical laws exists
 - If only a few parameters must be learned

Does Version 0.5 work?

(in general) **NO!**



1. Execute initial policy $\pi_0(\mathbf{u}_t|\mathbf{x}_t)$ (i.e. a random policy) and collect data $\mathcal{D} = \{(\mathbf{x}, \mathbf{u}, \mathbf{x}')_i\}$
2. Learn dynamics $f(\mathbf{x}, \mathbf{u})$ that minimizes $\sum_i \|f(\mathbf{x}_i, \mathbf{u}_i) - \mathbf{x}'_i\|^2$
3. Backpropagate $f(\mathbf{x}, \mathbf{u})$ and calculate sequence of actions (i.e. iLQR) $\rightarrow \pi_f(\mathbf{u}_t|\mathbf{x}_t)$

$$p_{\pi_0}(\mathbf{x}_t) \neq p_{\pi_f}(\mathbf{x}_t)$$

(Distribution Mismatch Problem)



Distribution Mismatch Problem increases if expressive classes of models are used (i.e. neural networks)

Can we do better?

Can we make $p_{\pi_0}(\mathbf{x}_t) = p_{\pi_f}(\mathbf{x}_t)$?




Need to collect data from $p_{\pi_f}(\mathbf{x}_t)$!


Modellbasiertes Reinforcement Learning Version 1.0

1. Execute initial policy $\pi_0(\mathbf{u}_t|\mathbf{x}_t)$ (i.e. a random policy) and collect data $\mathcal{D} = \{(\mathbf{x}, \mathbf{u}, \mathbf{x}')_i\}$
2. Learn dynamics $f(\mathbf{x}, \mathbf{u})$ that minimizes $\sum_i \|f(\mathbf{x}_i, \mathbf{u}_i) - \mathbf{x}'_i\|^2$
3. Backpropagate $f(\mathbf{x}, \mathbf{u})$ and calculate sequence of actions (i.e. iLQR)
4. Execute those actions and add the resulting data $\{(\mathbf{x}, \mathbf{u}, \mathbf{x}')_i\}$ to \mathcal{D}

What happens if the dynamic models contains little error?



Can we do better?



**Model predictive control
helps with error**

Modellbasiertes Reinforcement Learning Version 1.5

1. Execute initial policy $\pi_0(\mathbf{u}_t|\mathbf{x}_t)$ (i.e. a random policy) and collect data $\mathcal{D} = \{(\mathbf{x}, \mathbf{u}, \mathbf{x}')_i\}$
 2. Learn dynamics $f(\mathbf{x}, \mathbf{u})$ that minimizes $\sum_i \|f(\mathbf{x}_i, \mathbf{u}_i) - \mathbf{x}'_i\|^2$
 3. Backpropagate $f(\mathbf{x}, \mathbf{u})$ and calculate sequence of actions (i.e. iLQR)
 4. Execute the first planned action, observe resulting state \mathbf{x}' (MPC)
 5. Append $(\mathbf{x}, \mathbf{u}, \mathbf{x}')$ to dataset \mathcal{D}
- every N steps

Summary

- Version 0.5: collect random samples, train dynamics, plan
 - Pro: simple, no iterative procedure
 - Con: distribution mismatch problem
- Version 1.0: iteratively collect data, replan, collect data
 - Pro: simple, solves distribution mismatch
 - Con: open loop plan might perform poorly, exp. in stochastic domains
- Version 1.5: iteratively collect data using MPC (replan in each step)
 - Pro: robust to small model errors
 - Con: computationally expensive, but have planning algorithm available

I. Modelbased Reinforcement Learning

- I. Learning of dynamic models
- II. Learning of dynamic models and policies


II. Representing a dynamic model

III. Global and local dynamic model

IV. Learning with local dynamic models with „Trust Regions“

What kind of models can we use?

Gaussian process



GP with input (\mathbf{x}, \mathbf{u}) and output \mathbf{x}'

Pro: very data-efficient

Con: not great with non-smooth dynamics

Con: very slow when dataset is big

Neural Network





image: Punjani & Abbeel '14

Input is (\mathbf{x}, \mathbf{u}) , output is \mathbf{x}'

Pro: very expressive, can use lots of data

Con: not so great in low data regimes

Gaussian Mixture Model




GMM over $(\mathbf{x}, \mathbf{u}, \mathbf{x}')$ tuples

Train on $(\mathbf{x}, \mathbf{u}, \mathbf{x}')$, condition to get $p(\mathbf{x}'|\mathbf{x}, \mathbf{u})$

For i 'th mixture element, $p_i(\mathbf{x}, \mathbf{u})$ gives region where the mode $p_i(\mathbf{x}'|\mathbf{x}, \mathbf{u})$ holds

Pro: very expressive, if the dynamics can be assumed as piecewise linear

Representation of Dynamic Models



I. Modelbased Reinforcement Learning

- I. Learning of dynamic models
- II. Learning of dynamic models and policies

II. Representing a dynamic model


III. Global and local dynamic model

IV. Learning with local dynamic models with „Trust Regions“

Challenges

Example: Global dynamic model $f(\mathbf{x}_t, \mathbf{u}_t)$ is represented by a neural network

Modellbasiertes Reinforcement Learning Version 1.0

- 
1. Execute initial policy $\pi_0(\mathbf{u}_t | \mathbf{x}_t)$ (i.e. a random policy) and collect data $\mathcal{D} = \{(\mathbf{x}, \mathbf{u}, \mathbf{x}')_i\}$
 2. Learn dynamics $f(\mathbf{x}, \mathbf{u})$ that minimizes $\sum_i \|f(\mathbf{x}_i, \mathbf{u}_i) - \mathbf{x}'_i\|^2$
 3. Backpropagate $f(\mathbf{x}, \mathbf{u})$ and calculate sequence of actions (i.e. iLQR)
 4. Execute those actions and add the resulting data $\{(\mathbf{x}, \mathbf{u}, \mathbf{x}')_i\}$ to \mathcal{D}

- Planner will seek out regions where the model is erroneously optimistic
- Need to find a very good model in most of the state space to converge on a good solution

The trouble with global models

- Planner will seek out regions where the model is erroneously optimistic
- Need to find a very good model in most of the state space to converge on a good solution
- In some tasks, the model is much more complex than the policy



Motivation

$$\min_{\mathbf{u}_1, \dots, \mathbf{u}_T} \sum_{t=1}^T c(\mathbf{x}_t, \mathbf{u}_t) \quad \text{s.t.} \quad \mathbf{x}_t = f(\mathbf{x}_{t-1}, \mathbf{u}_{t-1})$$



$$\min_{\mathbf{u}_1, \dots, \mathbf{u}_T} c(\mathbf{x}_1, \mathbf{u}_1) + c(f(\mathbf{x}_1, \mathbf{u}_1), \mathbf{u}_2) + \dots + c(f(f(\dots) \dots), \mathbf{u}_T)$$

Usual story: differentiate via backpropagation and optimize (i.e. iLQR)

need: $\frac{\partial f}{\partial \mathbf{x}_t}, \frac{\partial f}{\partial \mathbf{u}_t}, \frac{\partial c}{\partial \mathbf{x}_t}, \frac{\partial c}{\partial \mathbf{u}_t}$


Approach

need:

$$\frac{\partial f}{\partial \mathbf{x}_t}, \frac{\partial f}{\partial \mathbf{u}_t}, \frac{\partial c}{\partial \mathbf{x}_t}, \frac{\partial c}{\partial \mathbf{u}_t}$$

idea: just fit $\frac{\partial f}{\partial \mathbf{x}_t}, \frac{\partial f}{\partial \mathbf{u}_t}$ around current trajectory or policy

$p(\mathbf{u}_t | \mathbf{x}_t)$ – time-varying linear-Gaussian controller –
can **execute** on the robot and produces
trajectory distribution




Learning a policy

$$p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t) = \mathcal{N}(f(\mathbf{x}_t, \mathbf{u}_t), \Sigma)$$

$$f(\mathbf{x}_t, \mathbf{u}_t) \approx \mathbf{A}_t \mathbf{x}_t + \mathbf{B}_t \mathbf{u}_t$$

$$\mathbf{A}_t = \frac{\partial f}{\partial \mathbf{x}_t} \quad \mathbf{B}_t = \frac{\partial f}{\partial \mathbf{u}_t}$$



Timedependent Linear Gaussian Controller

run $p(\mathbf{u}_t|\mathbf{x}_t)$ on robot,
collect $\mathcal{D} = \{\tau_i\}$



iLQR produces: $\hat{\mathbf{x}}_t, \hat{\mathbf{u}}_t, \mathbf{K}_t, \mathbf{k}_t$

$$p(\mathbf{u}_t|\mathbf{x}_t) = \mathcal{N}(\mathbf{K}_t(\mathbf{x}_t - \hat{\mathbf{x}}_t) + \mathbf{k}_t + \hat{\mathbf{u}}_t, \Sigma_t)$$

Set $\Sigma_t = \mathbf{Q}_{\mathbf{u}_t, \mathbf{u}_t}^{-1}$


$Q(\mathbf{x}_t, \mathbf{u}_t)$ is the cost to go: total cost we get after taking an action \mathbf{u}_t

$$Q(\mathbf{x}_t, \mathbf{u}_t) = \text{const} + \frac{1}{2} \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{Q}_t \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix} + \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{q}_t$$

$\mathbf{Q}_{\mathbf{u}_t, \mathbf{u}_t}$ is big, if changing \mathbf{u}_t
changes the Q-value a lot!



If \mathbf{u}_t changes Q-value a lot, don't
vary \mathbf{u}_t so much. Exploration
noise Σ_t must be low



Time-dependent Linear Gaussian Controller

run $p(\mathbf{u}_t|\mathbf{x}_t)$ on robot,
collect $\mathcal{D} = \{\tau_i\}$



iLQR produces: $\hat{\mathbf{x}}_t, \hat{\mathbf{u}}_t, \mathbf{K}_t, \mathbf{k}_t$

$$p(\mathbf{u}_t|\mathbf{x}_t) = \mathcal{N}(\mathbf{K}_t(\mathbf{x}_t - \hat{\mathbf{x}}_t) + \mathbf{k}_t + \hat{\mathbf{u}}_t, \Sigma_t)$$


Set $\Sigma_t = \mathbf{Q}_{\mathbf{u}_t, \mathbf{u}_t}^{-1}$

Standard LQR solves $\min_{\mathbf{u}_1, \dots, \mathbf{u}_T} \sum_{t=1}^T c(\mathbf{x}_t, \mathbf{u}_t)$

Linear-Gaussian solution solves $\min_{\mathbf{u}_1, \dots, \mathbf{u}_T} \sum_{t=1}^T E[c(\mathbf{x}_t, \mathbf{u}_t) - \mathcal{H}(p(\mathbf{u}_t|\mathbf{x}_t))]$

Maximum Entropy: act as randomly as possible while minimizing cost


– Entropy: A measure for the average information content



Linearize local dynamics

Linearize local dynamic models

$$p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t)$$



$$\{(\mathbf{x}_t, \mathbf{u}_t, \mathbf{x}_{t+1})_i\}$$


Version 1.0: Linearize $p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t)$ at each time step using linear regression

$$p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t) = \mathcal{N}(\mathbf{A}_t\mathbf{x}_t + \mathbf{B}_t\mathbf{u}_t + \mathbf{c}_t, \mathbf{N}_t) \quad \mathbf{A}_t \approx \frac{\partial f}{\partial \mathbf{x}_t} \quad \mathbf{B}_t \approx \frac{\partial f}{\partial \mathbf{u}_t}$$

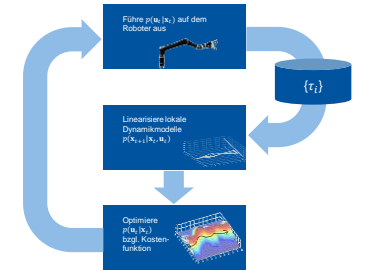

Can we do better?

Version 2.0: Linearize $p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t)$ using *Bayesian* linear regression



- Bayesian linear regression uses prior: $p(\mathbf{x}_t, \mathbf{u}_t)p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t) = p(\mathbf{x}_t, \mathbf{u}_t, \mathbf{x}_{t+1})$
- Use your favourite global model as a prior (GP, deep net, GMM)



How to stay close to old controller?



How to stay close to old controller?



$$p(\mathbf{u}_t|\mathbf{x}_t) = \mathcal{N}(\mathbf{K}_t(\mathbf{x}_t - \hat{\mathbf{x}}_t) + \mathbf{k}_t + \hat{\mathbf{u}}_t, \Sigma_t)$$

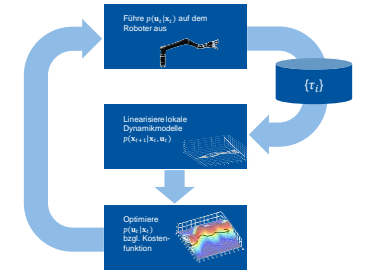
$$p(\tau) = p(\mathbf{x}_1) \prod_{t=1}^T p(\mathbf{u}_t|\mathbf{x}_t) p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t)$$

New trajectory distribution $p(\tau)$ must be similar to the old one $\bar{p}(\tau)$

If trajectory distribution is close, the dynamics will be close too!

What does “close” mean?

Kullback-Leibler divergence: $D_{KL}(p(\tau)||\bar{p}(\tau)) < \varepsilon$ **From here comes a lot of mathematics!**



Its your turn!




Visit the website and implement it!



Introduction to the tasks

Tasks for today and tomorrow

- Task 1:
 - Implement an LQR Backward and Forward pass
 - Try to understand it!
 - Test it with our test method
- Task 2:
 - Implement linearization of the dynamic model
 - Try to understand it!
 - Test it with our test-method
 - Test it on the Box2D Scenario
- Task 3:
 - Test it with Kinova Jaco 2 in simulation
 - Adjust cost function
- Task 4:
 - Test it with real Kinova Jaco 2
 - Adjust cost function



Task 1 – Installation procedure

Download source code (do it in your home directory: `cd ~`):

```
git clone https://github.com/philippente/ss2017\_task1\_lqr.git
```

Edit `.bashrc` to set environment variables:

```
gedit ~/.bashrc
```

At the end of file, the lines should look like this:

```
source /opt/ros/indigo/setup.bash
source /home/useradmin/catkin_ws/devel/setup.bash
export
ROS_PACKAGE_PATH=$ROS_PACKAGE_PATH:/opt/ros/indigo/share:/opt/ros/indigo/stacks:/home/useradmin/ss2017_task1_lqr:/home/useradmin/ss2017_task1_lqr/src/gps_agent_pkg
```

Check if the blue part of the source folder and `ROS_PACKAGE_PATH` is correct!

Then save it and close it. Source the `.bashrc` (load the environment variables):

```
source ~/.bashrc
```

Now, compile some stuff:

```
cd ss2017_task1_lqr
sh compile_proto.sh
cd /src/gps_agent_pkg
cmake .
make -j
```

Task 1 – Installation procedure

- Open PyCharm
- Import the folder *ss2017_task1_lqr* as a new project
- Open within PyCharm: `python/gps/algorithm/algorithm_traj_opt.py`
- **Task: Implement the forward and backward pass of an LQR! Look at the website for advices` : <https://goo.gl/X5twgi>**
- You can test your implementation with a little test program
 - using a terminal, open the directory *ss2017_task1_lqr*
 - Start the program with: `python python/gps/lqr_test.py`
 - Was it successful?

Thanks for your attention!

- Univariate Gaussian
 - Multivariate Gaussian
 - Law of Total Probability
 - Conditioning (Bayes' rule)
-
- Disclaimer: lots of linear algebra in next few lectures. In fact, pretty much all computations with Gaussians will be reduced to linear algebra!