




Source: <https://blenderartists.org/forum/showthread.php?316399-Industrial-Robots>

# Robotics for Future Industrial Applications

## Summary Week 2 – Intelligent Motion Planning

Philipp Ennen, M.Sc.

# Organizational

Time	Sunday 13 <sup>th</sup> August	Monday 14 <sup>th</sup> August	Tuesday 15 <sup>th</sup> August	Wednesday 16 <sup>th</sup> August	Thursday 17 <sup>th</sup> August	Friday 18 <sup>th</sup> August	Saturday 19 <sup>th</sup> August	
9:00-9:30	Free Time for Excursions, Sight-Seeing and Self-Study	Introduction to Industrial Robots & Challenges	Fundamentals of Robot Learning and Control Theory		Robot Learning with iterative Linear- Quadratic Regulator	Practical Unit: Robot Learning of an Assembly Task	Free Time for Excursions, Sight-Seeing and Self-Study	
9:30-10:00								
10:00-10:30								
10:30-11:00								
11:00-11:30		Motion Planning	Demonstration			Practical Unit: Robot Learning of an Assembly Task		
11:30-12:00								
12:00-12:30		Lunch Break at “Bistro Kerres”			Lunch Break at “Bistro Kerres”			
12:30-13:00								
13:00-13:30								
13:30-14:00								
14:00-14:30		Practical Unit: Introduction to Linux and Robot Operating System	Programming using Python		Practical Unit: Robot Learning of an Assembly Task	Show: Let the robot learn!		
14:30-15:00								
15:00-15:30			Practical Unit: Programming using Python					
15:30-16:00								
16:00-16:30								
16:30-17:00								
17:00-17:30								
17:30-18:00								
18:00-18:30								
18:30-19:00								
19:00-19:30								
19:30-20:00								
20:00-20:30				OPTIONAL International Tuesday with the INCAS Student Organisation 		„Karaoke Night“ with the Summer School Team and your Buddies 		
20:30-21:00								
21:00-21:30								
21:30-22:00								



# Robots as an Example for Intelligent Machines

## The central questions of robotics

Three main question:

Where am I?



- Easy in Industrial Robotics
- Hard in Mobile Robotics


Where should I go?



- Often hand-engineered
- Or a Result of Task Planning

This weeks focus!

How do I go there?

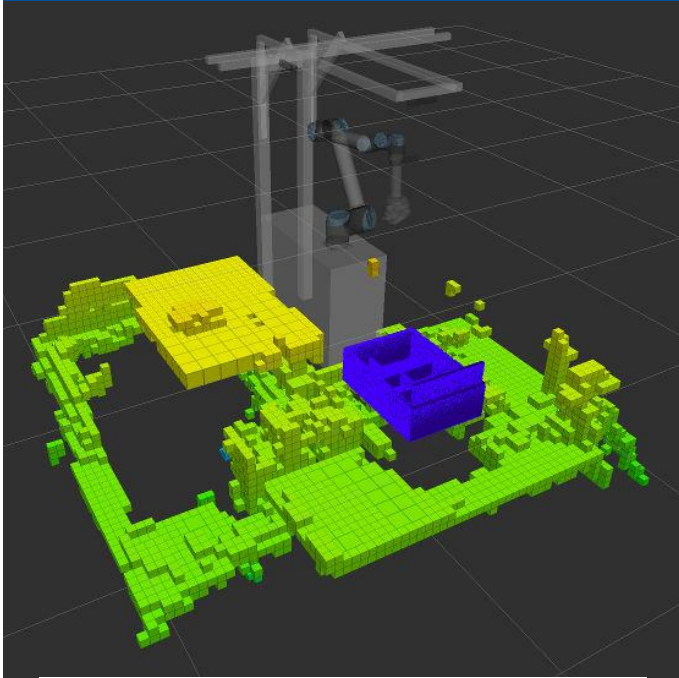


- Comparable Easy in Mobile Robotics
- Hard in Industrial Robotics


# Robots as an Example for Intelligent Machines

## How do I (the robot) go there?

### Model of the Robot and Environment



### Sensing




This weeks topic!

### Motion Planning

Requires Goalstate:

- i.e. hand-engineered
- i.e. via a cost function

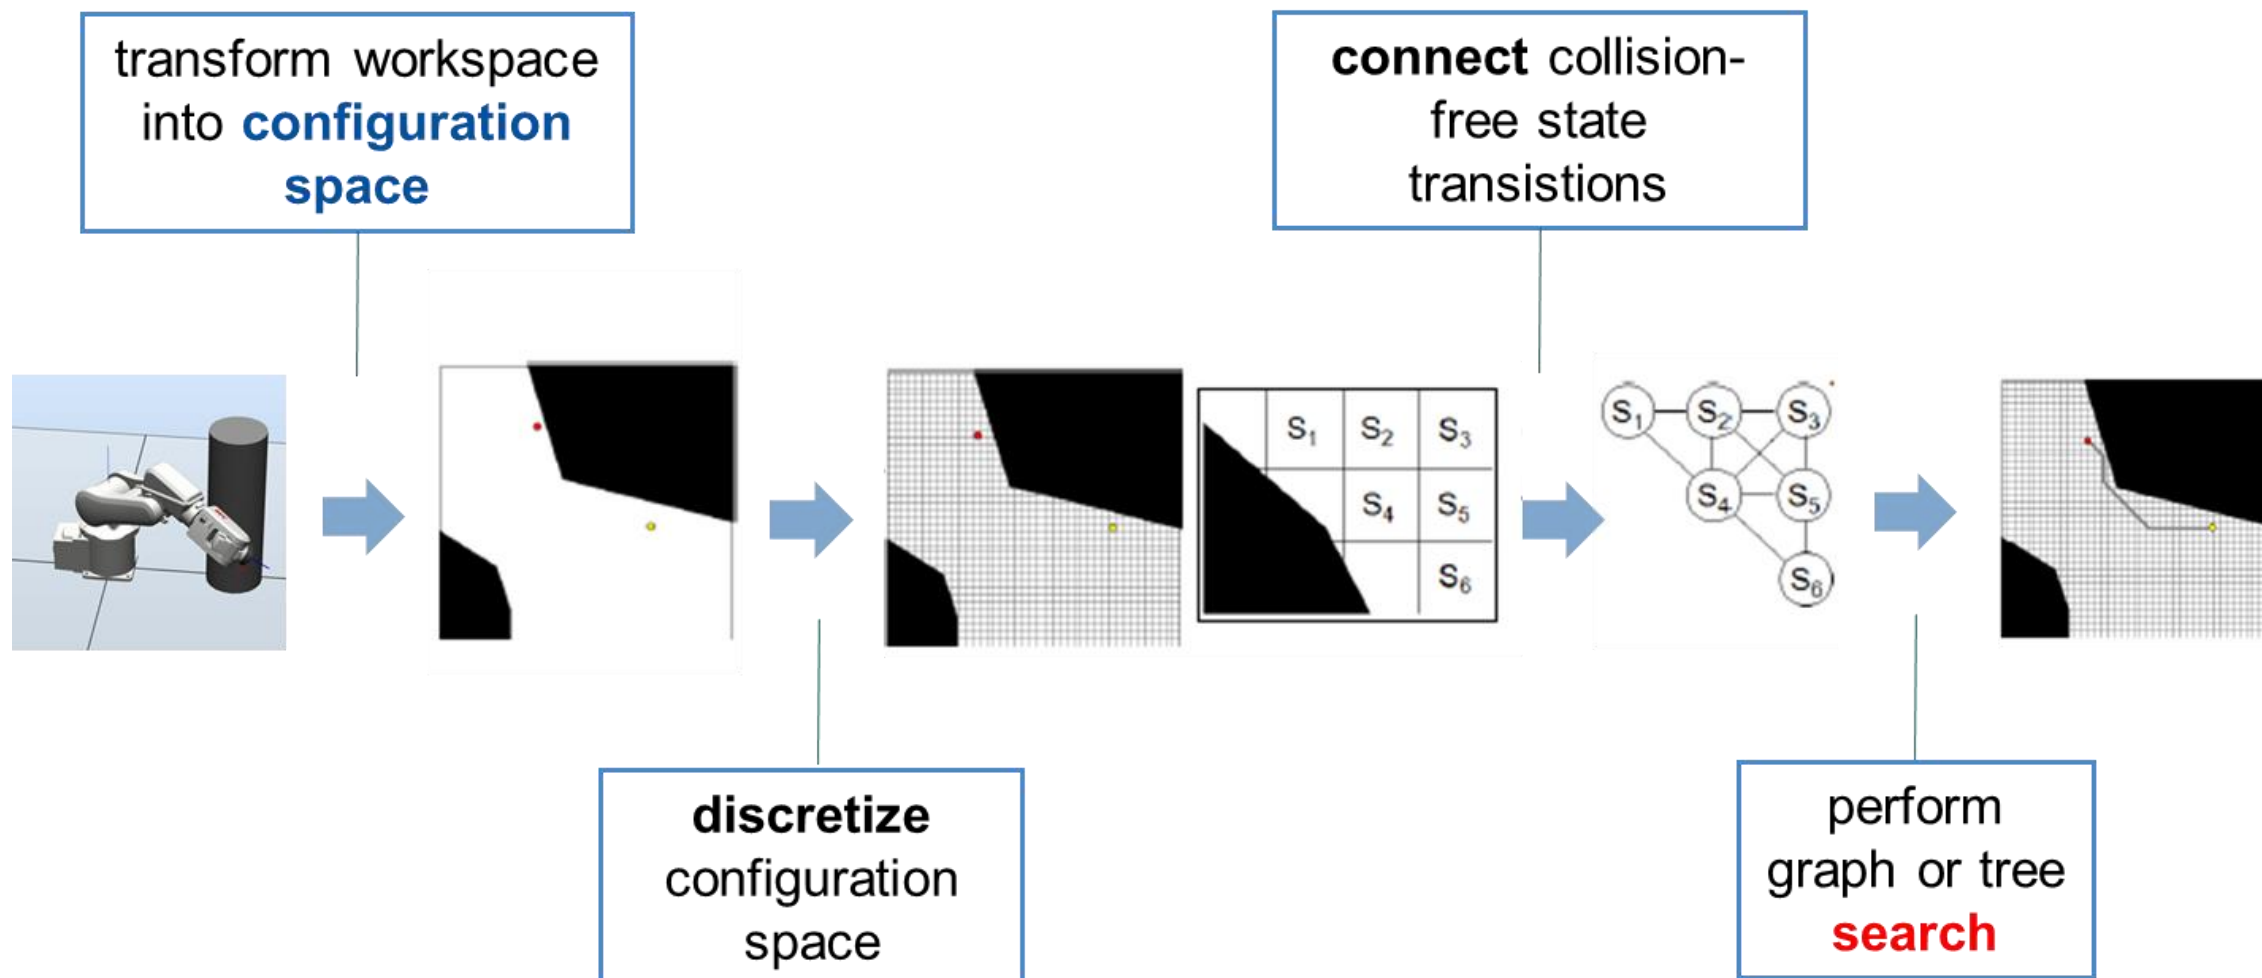
### Motion Control and Execution



$$\begin{bmatrix} \dot{x} \\ \dot{y} \\ \dot{\phi} \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \frac{-(I+ml^2)b}{I(M+m)+Mml^2} & \frac{m^2gl^2}{I(M+m)+Mml^2} \\ 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ \phi \end{bmatrix} + \begin{bmatrix} 0 \\ \frac{I+ml^2}{I(M+m)+Mml^2} \\ \frac{mgl(M+m)}{I(M+m)+Mml^2} \end{bmatrix} u$$

## Motion Planning: Configuration Space

### Motion Planning Pipeline within $\mathcal{C}_{space}$

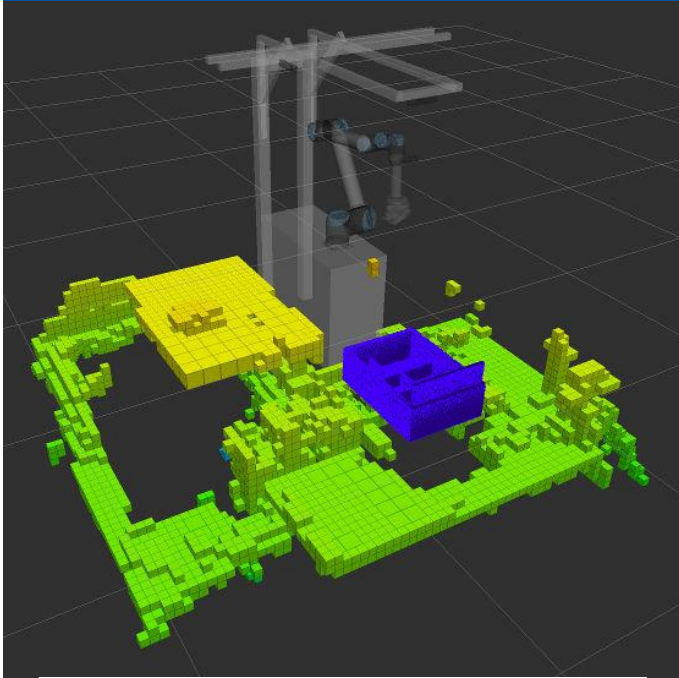




# Robots as an Example for Intelligent Machines


## How do I (the robot) go there?

### Model of the Robot and Environment



$$\begin{bmatrix} \ddot{x} \\ \ddot{y} \\ \ddot{\phi} \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \frac{-(I + ml^2)b}{I(M+m) + Mml^2} & \frac{m^2 gl^2}{I(M+m) + Mml^2} \\ 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ \phi \end{bmatrix} + \begin{bmatrix} 0 \\ \frac{I + ml^2}{I(M+m) + Mml^2} \\ \frac{mgl(M+m)}{I(M+m) + Mml^2} \end{bmatrix} u$$

### Sensing



What if the model of the robot and environment is hard to describe (or unknown)?


Think about flexible objects!




Requires Goalstate:

- i.e. hand-engineered
- i.e. via a cost function

### Motion Control and Execution




Think about contact-situations!



## Robots as an Example for Intelligent Machines

What if the model of the robot and environment is hard to describe (or unknown)?

- Use an Reinforcement Learning Agent!




# Robots as an Example for Intelligent Machines

What if the model of the robot and environment is hard to describe (or unknown)?

This weeks topic!

## Model-based RL:

- Learn to predict next state:  $P(s'|s, a)$
- Learn to predict immediate reward  $P(r'|s, a)$



## Model-free RL:

- Learn to predict value:  $V(s)$  or  $Q(s, a)$



## Linear Quadratic Regulator

- Special case: Systems with
  - **Linear dynamics**
  - **Quadratic costs**
- LQR provides an exact solution

$$\min_{\mathbf{u}_1, \dots, \mathbf{u}_T} c(\mathbf{x}_1, \mathbf{u}_1) + c(f(\mathbf{x}_1, \mathbf{u}_1), \mathbf{u}_2) + \dots + c(f(f(\dots)), \mathbf{u}_T)$$

$$f(\mathbf{x}_t, \mathbf{u}_t) = \mathbf{F}_t \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix} + \mathbf{f}_t$$

linear

$$c(\mathbf{x}_t, \mathbf{u}_t) = \frac{1}{2} \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{C}_t \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix} + \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{c}_t$$

quadratic

# Linear Quadratic Regulator

## Pseudocode Algorithm

Backward recursion  
(Get linear equations for  $\mathbf{u}$ )

for  $t = T$  to 1:

$$\mathbf{Q}_t = \mathbf{C}_t + \mathbf{F}_t^T \mathbf{V}_{t+1} \mathbf{F}_t$$

$$\mathbf{q}_t = \mathbf{c}_t + \mathbf{F}_t^T \mathbf{V}_{t+1} \mathbf{f}_t + \mathbf{F}_t^T \mathbf{v}_{t+1}$$

$$Q(\mathbf{x}_t, \mathbf{u}_t) = \text{const} + \frac{1}{2} \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{Q}_t \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix} + \begin{bmatrix} \mathbf{x}_t \\ \mathbf{u}_t \end{bmatrix}^T \mathbf{q}_t$$

$$\mathbf{u}_t \leftarrow \arg \min_{\mathbf{u}_t} Q(\mathbf{x}_t, \mathbf{u}_t) = \mathbf{K}_t \mathbf{x}_t + \mathbf{k}_t$$


$$\mathbf{K}_t = -\mathbf{Q}_{\mathbf{u}_t, \mathbf{u}_t}^{-1} \mathbf{Q}_{\mathbf{u}_t, \mathbf{x}_t}$$

$$\mathbf{k}_t = -\mathbf{Q}_{\mathbf{u}_t, \mathbf{u}_t}^{-1} \mathbf{q}_{\mathbf{u}_t}$$

$$\mathbf{V}_t = \mathbf{Q}_{\mathbf{x}_t, \mathbf{x}_t} + \mathbf{Q}_{\mathbf{x}_t, \mathbf{u}_t} \mathbf{K}_t + \mathbf{K}_t^T \mathbf{Q}_{\mathbf{u}_t, \mathbf{x}_t} \mathbf{x}_t + \mathbf{K}_t^T \mathbf{Q}_{\mathbf{u}_t, \mathbf{u}_t} \mathbf{K}_t$$

$$\mathbf{v}_t = \mathbf{K}_t^T \mathbf{Q}_{\mathbf{u}_t, \mathbf{u}_t} \mathbf{k}_t + \mathbf{Q}_{\mathbf{x}_t, \mathbf{u}_t} \mathbf{k}_t + \mathbf{q}_{\mathbf{x}_t} + \mathbf{K}_t^T \mathbf{q}_{\mathbf{u}_t}$$

$$V(\mathbf{x}_t) = \text{const} + \frac{1}{2} \mathbf{x}_t^T \mathbf{V}_t \mathbf{x}_t + \mathbf{x}_t^T \mathbf{v}_t$$



Forward recursion  
(Use known initial state to get  
values for  $\mathbf{u}$ )

for  $t = 1$  to  $T$ :

$$\mathbf{u}_t = \mathbf{K}_t \mathbf{x}_t + \mathbf{k}_t$$

$$\mathbf{x}_{t+1} = f(\mathbf{x}_t, \mathbf{u}_t)$$

### Why do we want to learn the dynamics?

- If  $\mathbf{x}_{t+1} = f(\mathbf{x}_t, \mathbf{u}_t)$  is known, we can do trajectory optimization
  - In the stochastic case  $p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t)$



Learn  $f(\mathbf{x}_t, \mathbf{u}_t)$  with subsequent backpropagation (i.e. iLQR)


### Modelbased Reinforcement Learning Version 0.5

1. Execute initial policy  $\pi_0(\mathbf{u}_t|\mathbf{x}_t)$  (i.e. a random policy) and collect data  $\mathcal{D} = \{(\mathbf{x}, \mathbf{u}, \mathbf{x}')_i\}$
2. Learn dynamics  $f(\mathbf{x}, \mathbf{u})$  that minimizes  $\sum_i \|f(\mathbf{x}_i, \mathbf{u}_i) - \mathbf{x}'_i\|^2$
3. Backpropagate  $f(\mathbf{x}, \mathbf{u})$  and calculate sequence of actions (i.e. iLQR)



## Does Version 0.5 work?

(in general) **NO!**



1. Execute initial policy  $\pi_0(\mathbf{u}_t|\mathbf{x}_t)$  (i.e. a random policy) and collect data  $\mathcal{D} = \{(\mathbf{x}, \mathbf{u}, \mathbf{x}')_i\}$
2. Learn dynamics  $f(\mathbf{x}, \mathbf{u})$  that minimizes  $\sum_i \|f(\mathbf{x}_i, \mathbf{u}_i) - \mathbf{x}'_i\|^2$
3. Backpropagate  $f(\mathbf{x}, \mathbf{u})$  and calculate sequence of actions (i.e. iLQR)  $\rightarrow \pi_f(\mathbf{u}_t|\mathbf{x}_t)$

$$p_{\pi_0}(\mathbf{x}_t) \neq p_{\pi_f}(\mathbf{x}_t)$$

(Distribution Mismatch Problem)



Distribution Mismatch Problem increases if expressive classes of models are used (i.e. neural networks)

Can we do better?

Can we make  $p_{\pi_0}(\mathbf{x}_t) = p_{\pi_f}(\mathbf{x}_t)$ ?



Need to collect data from  $p_{\pi_f}(\mathbf{x}_t)$ !


### Modellbasiertes Reinforcement Learning Version 1.0

1. Execute initial policy  $\pi_0(\mathbf{u}_t|\mathbf{x}_t)$  (i.e. a random policy) and collect data  $\mathcal{D} = \{(\mathbf{x}, \mathbf{u}, \mathbf{x}')_i\}$
2. Learn dynamics  $f(\mathbf{x}, \mathbf{u})$  that minimizes  $\sum_i \|f(\mathbf{x}_i, \mathbf{u}_i) - \mathbf{x}'_i\|^2$
3. Backpropagate  $f(\mathbf{x}, \mathbf{u})$  and calculate sequence of actions (i.e. iLQR)
4. Execute those actions and add the resulting data  $\{(\mathbf{x}, \mathbf{u}, \mathbf{x}')_i\}$  to  $\mathcal{D}$

$$p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t) = \mathcal{N}(f(\mathbf{x}_t, \mathbf{u}_t), \Sigma)$$

$$f(\mathbf{x}_t, \mathbf{u}_t) \approx \mathbf{A}_t \mathbf{x}_t + \mathbf{B}_t \mathbf{u}_t$$

$$\mathbf{A}_t = \frac{\partial f}{\partial \mathbf{x}_t} \quad \mathbf{B}_t = \frac{\partial f}{\partial \mathbf{u}_t}$$





### Linearized local dynamics

Goal: get the system dynamics  $p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t)$  for each timestep  $t$

Data: samples generated by the previous controller  $\hat{p}_i(\mathbf{u}_t|\mathbf{x}_t) \rightarrow \{(\mathbf{x}_t, \mathbf{u}_t, \mathbf{x}_{t+1})_i\}$


Linear Gaussian Dynamics are defined as

$$p(\mathbf{x}_{t+1}|\mathbf{x}_t, \mathbf{u}_t) = \mathcal{N}(f_{xt}\mathbf{x}_t + f_{ut}\mathbf{u}_t + f_{ct}, \mathbf{F}_t)$$

How can we determine linear  
Gaussian dynamics from few  
samples?

# What kind of models can we use?

## Gaussian process



GP with input  $(\mathbf{x}, \mathbf{u})$  and output  $\mathbf{x}'$

Pro: very data-efficient

Con: not great with non-smooth dynamics

Con: very slow when dataset is big

## Neural Network




image: Punjani & Abbeel '14


Input is  $(\mathbf{x}, \mathbf{u})$ , output is  $\mathbf{x}'$

Pro: very expressive, can use lots of data

Con: not so great in low data regimes

This week's focus!

## Gaussian Mixture Model



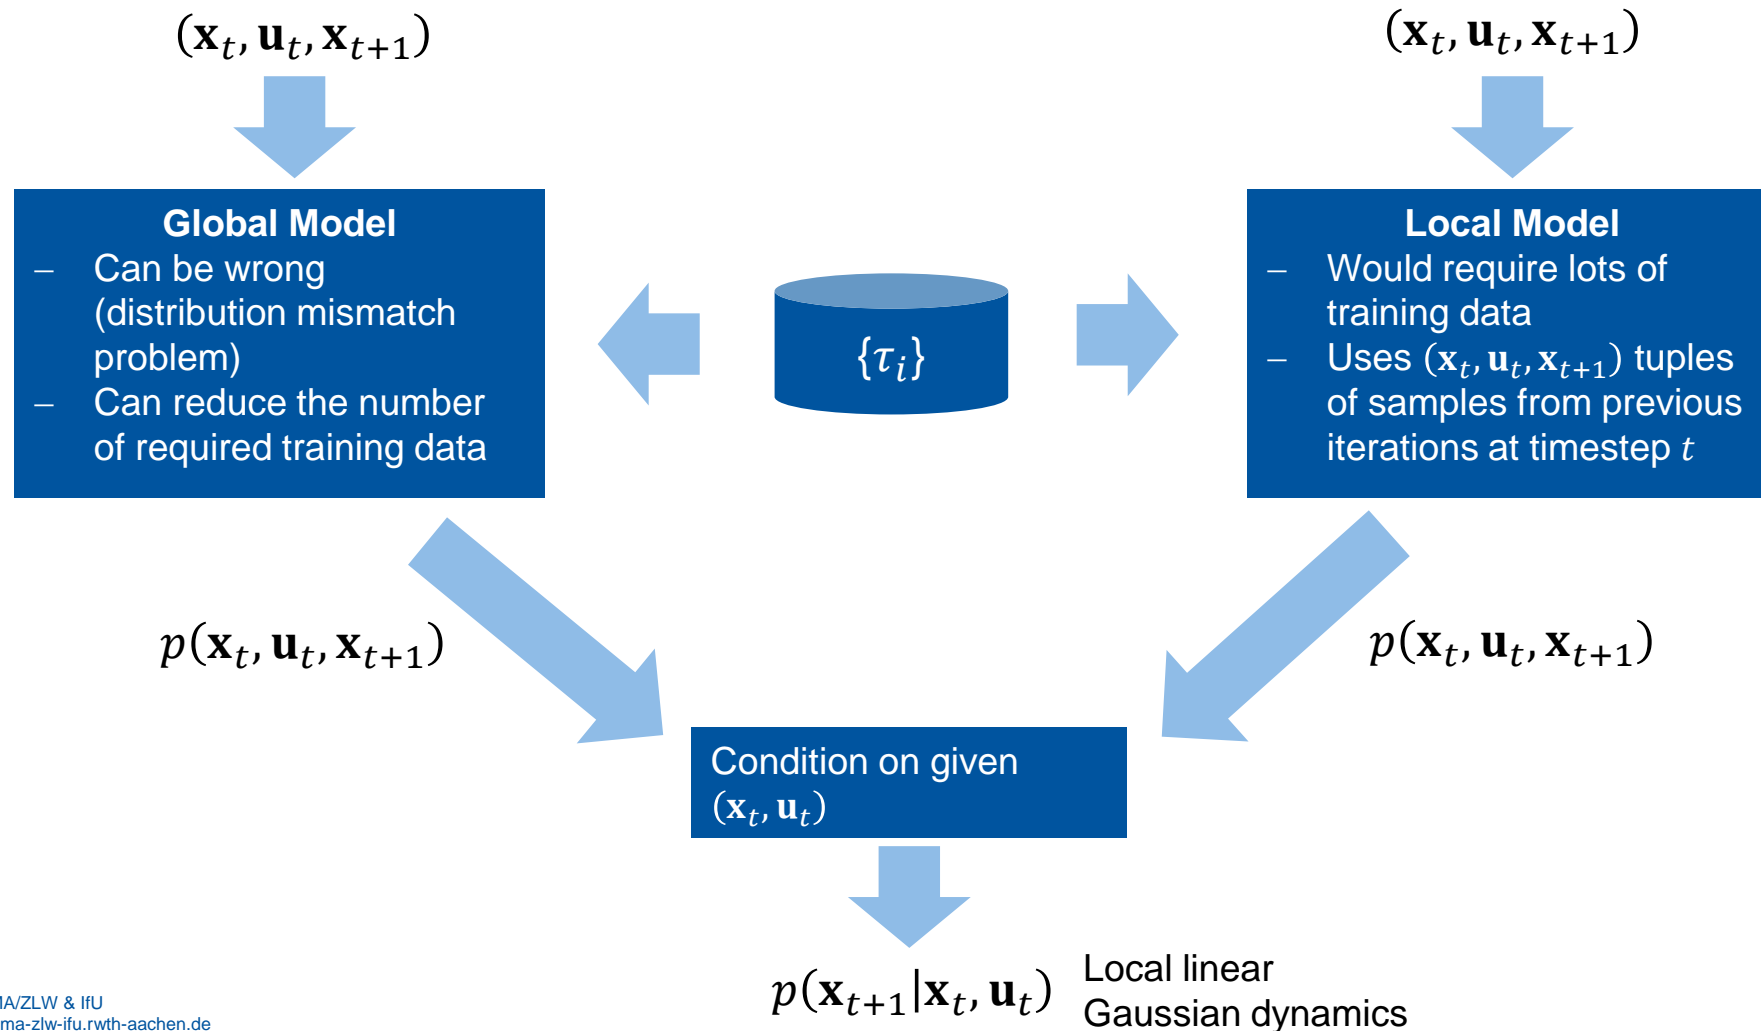
GMM over  $(\mathbf{x}, \mathbf{u}, \mathbf{x}')$  tuples

Train on  $(\mathbf{x}, \mathbf{u}, \mathbf{x}')$ , condition to get  $p(\mathbf{x}'|\mathbf{x}, \mathbf{u})$

For  $i$ 'th mixture element,  $p_i(\mathbf{x}, \mathbf{u})$  gives region where the mode  $p_i(\mathbf{x}'|\mathbf{x}, \mathbf{u})$  holds

Pro: very expressive, if the dynamics can be assumed as piecewise linear

## Combining global and local models





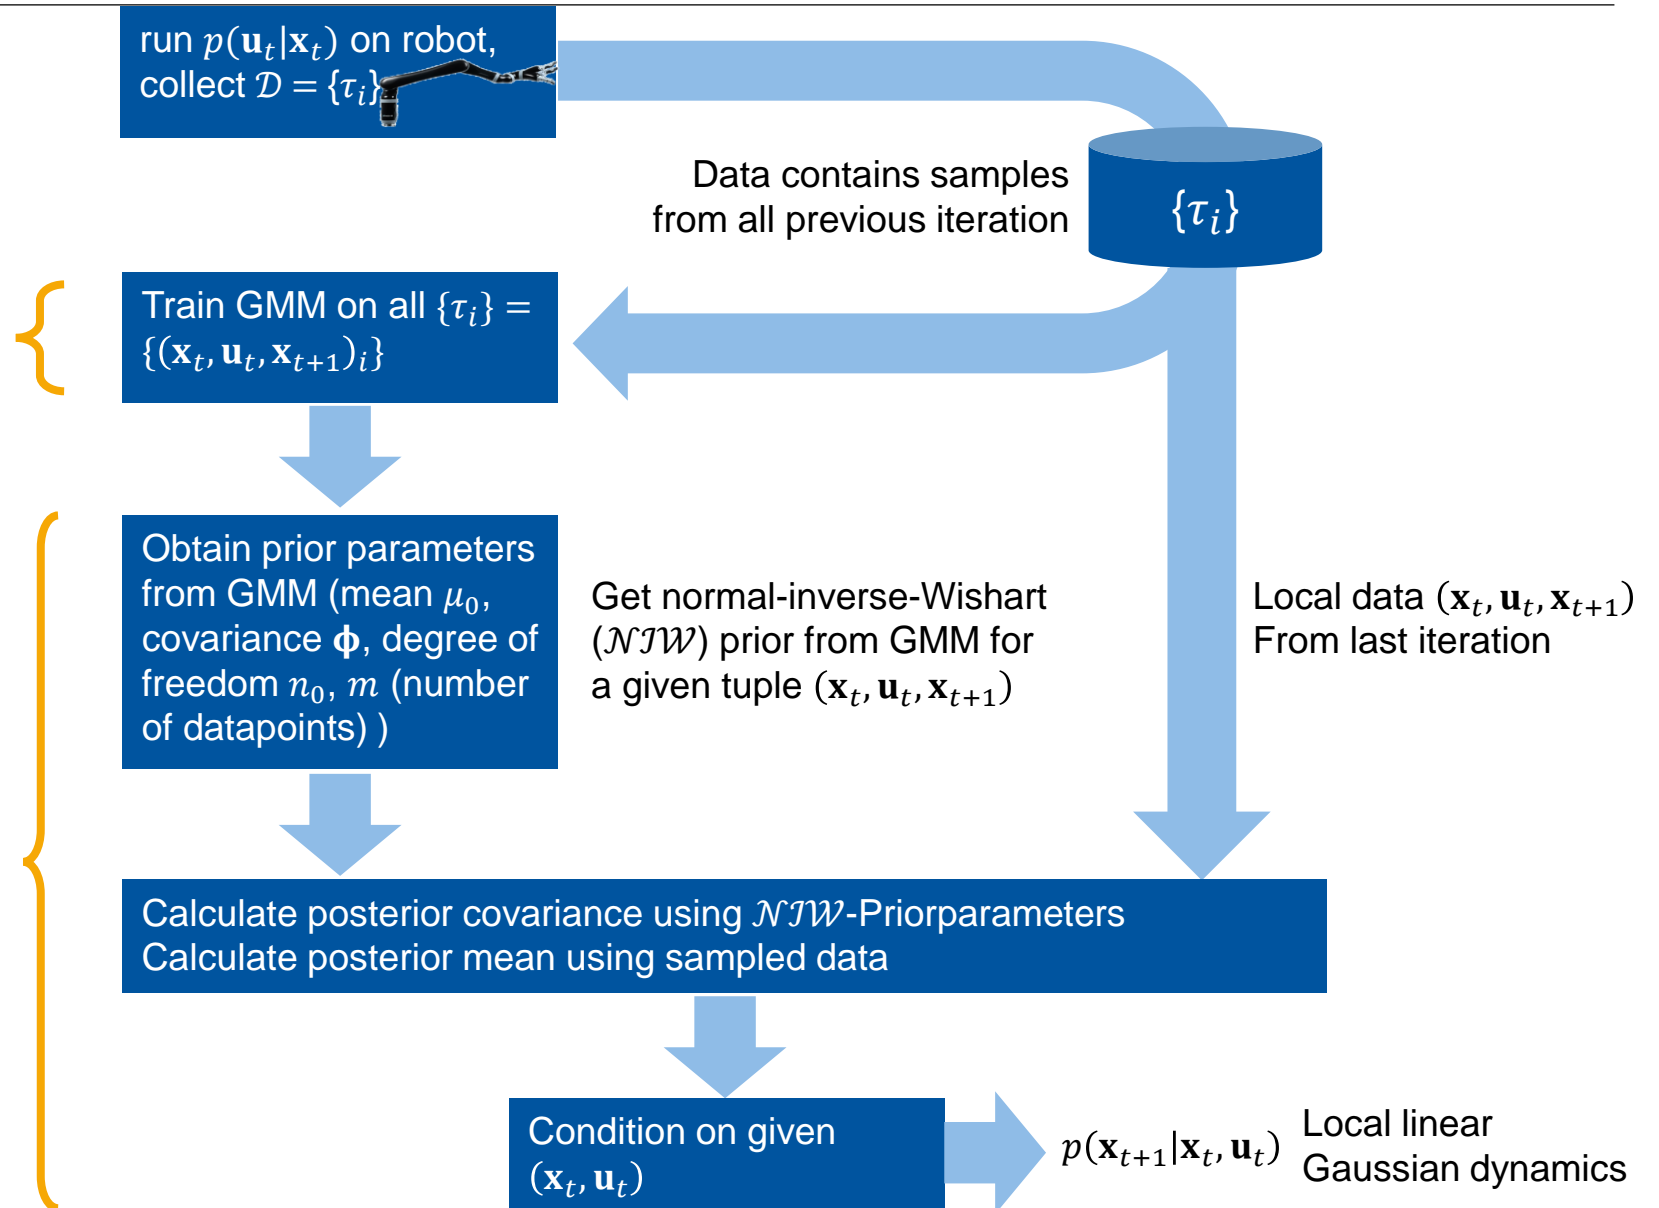
# Learning Local and Global Models

## Train GMM: **Global Dynamic Model**

- Uses data from nearby timesteps
- Uses data from prior iterations

## Linearize: **Local Dynamic Model**

- Uses prior from local dynamic model
- Uses data from last iteration at timestep  $t$
- Condition on given  $(\mathbf{x}_t, \mathbf{u}_t)$



**Thank you  
for your attention!**